

When Is Graph Reordering An Optimization?

A Cross Application and Input Graph Study on the Effectiveness of Lightweight Graph Reordering

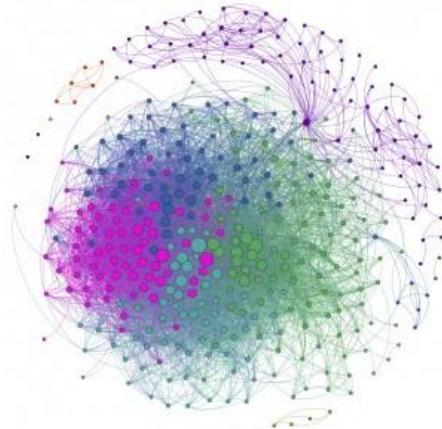
Vignesh Balaji

Brandon Lucia

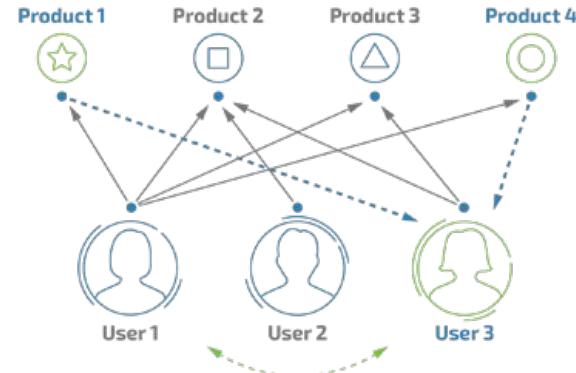
Graph Processing Has Many Applications



Path Planning



Social network analysis

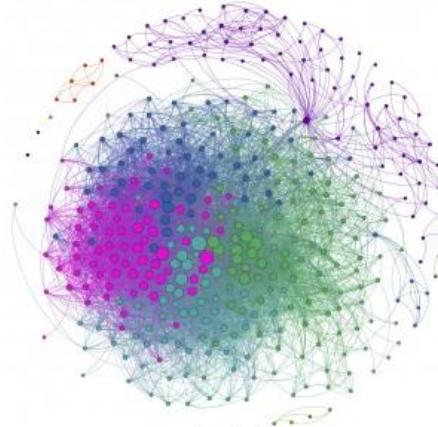


Recommender systems

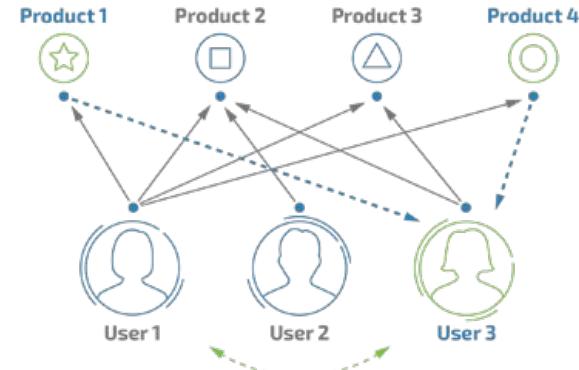
Graph Processing Has Many Applications



Path Planning



Social network analysis

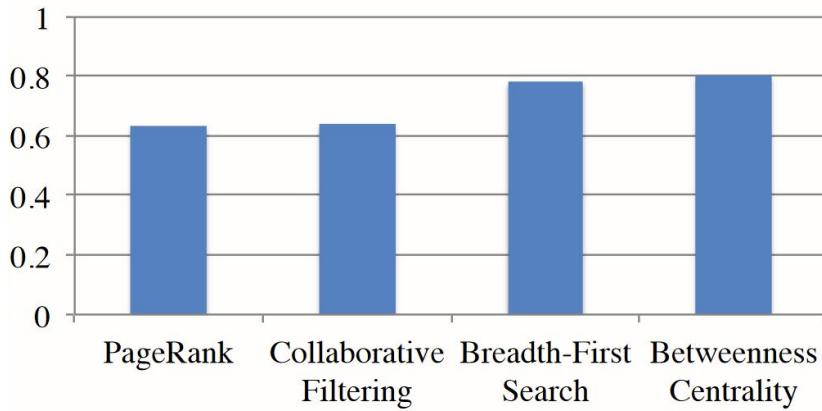


Recommender systems

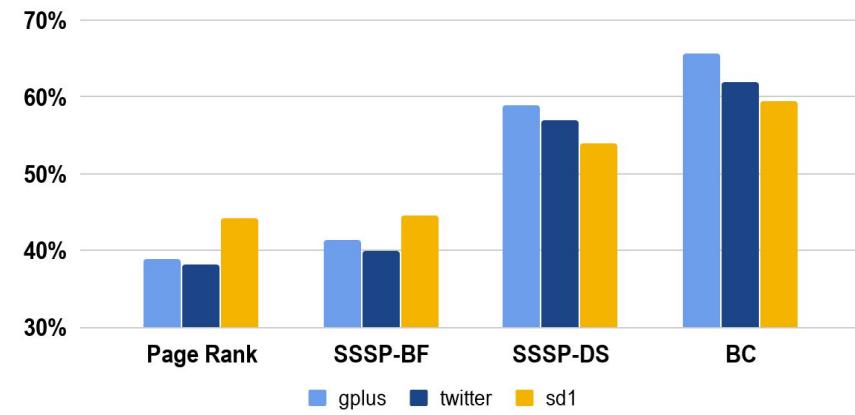


Graph Applications Are Memory Bound

Cycles stalled on DRAM / Total Cycles

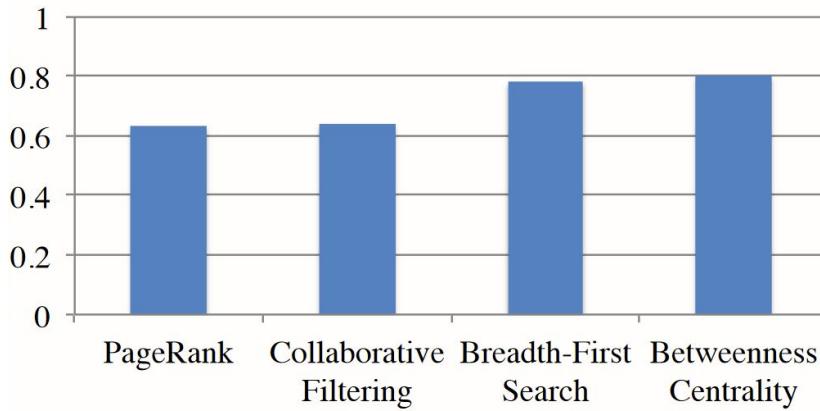


LLC Miss Rate

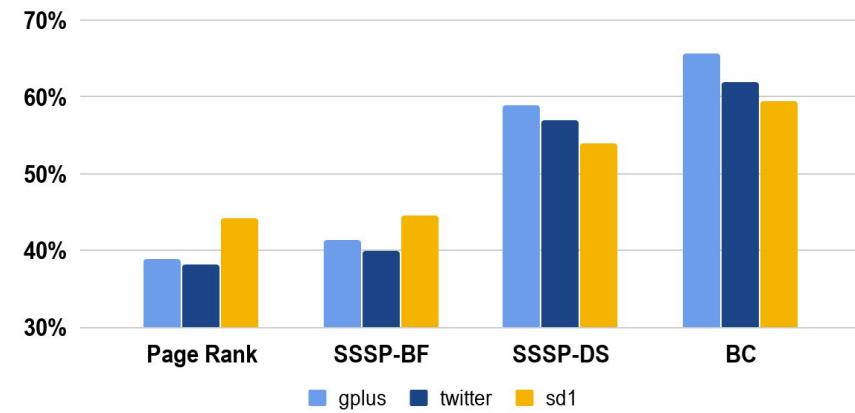


Graph Applications Are Memory Bound

Cycles stalled on DRAM / Total Cycles



LLC Miss Rate



Problem: Poor LLC locality \Rightarrow Many long-latency DRAM accesses

Reason For Poor Locality - Irregular Memory Accesses

Reason For Poor Locality - Irregular Memory Accesses

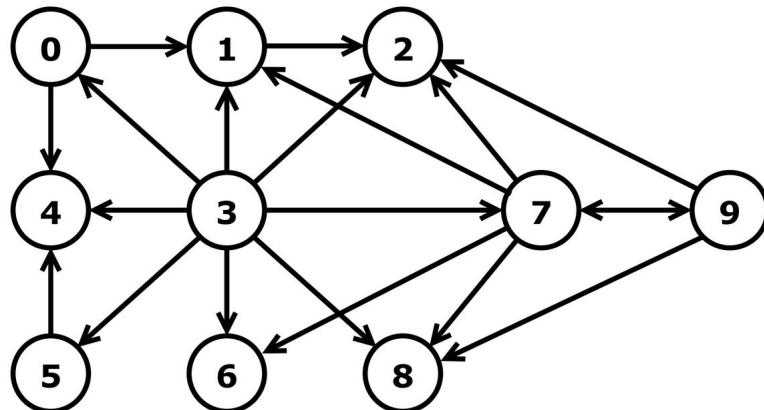
```
for v in G:  
    for u in neigh(v):  
        process(..., vtxData[u], ...)
```

Typical graph processing kernel

Reason For Poor Locality - Irregular Memory Accesses

```
for v in G:  
    for u in neigh(v):  
        process(..., vtxData[u], ...)
```

Typical graph processing kernel

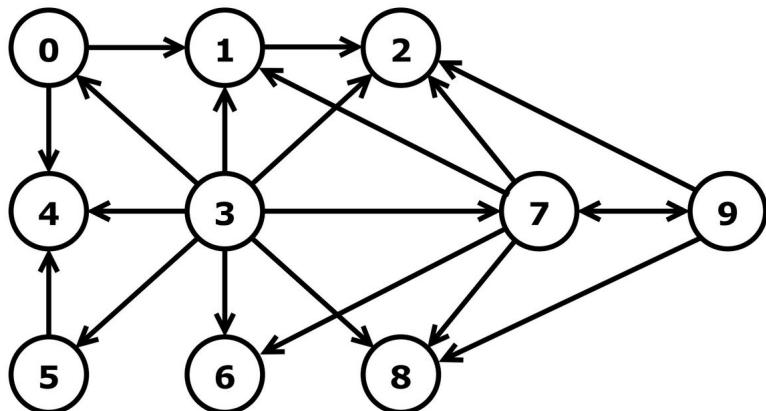


Input Graph

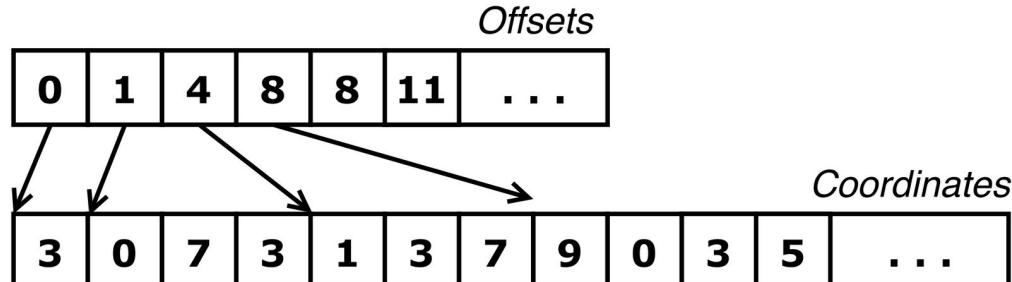
Reason For Poor Locality - Irregular Memory Accesses

```
for v in G:  
    for u in neigh(v):  
        process(..., vtxData[u], ...)
```

Typical graph processing kernel



Input Graph



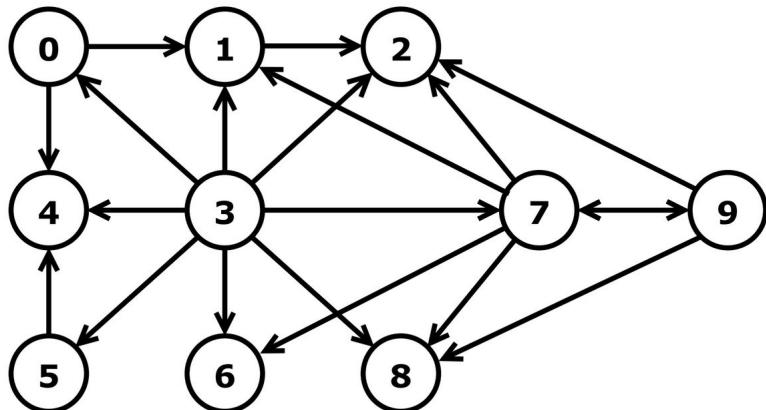
Compressed Sparse Row (CSR) Representation

Reason For Poor Locality - Irregular Memory Accesses

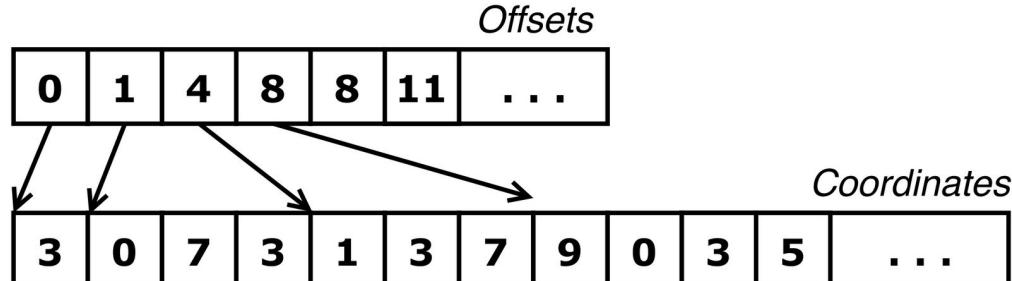
```
for v in G:  
    for u in neigh(v):  
        process(..., vtxData[u], ...)
```

Irregular accesses to
vtxData array

Typical graph processing kernel



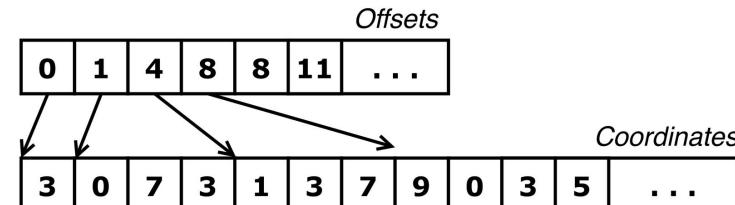
Input Graph



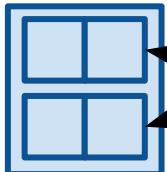
Compressed Sparse Row (CSR) Representation

Irregular Accesses Have Poor Temporal And Spatial Locality

```
for v in G:  
    for u in neigh(v):  
        process(..., vtxData[u], ...)
```



Time



2 lines

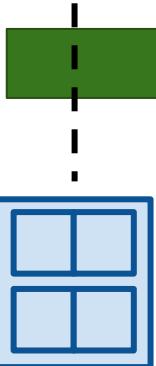
2 words/line

LLC

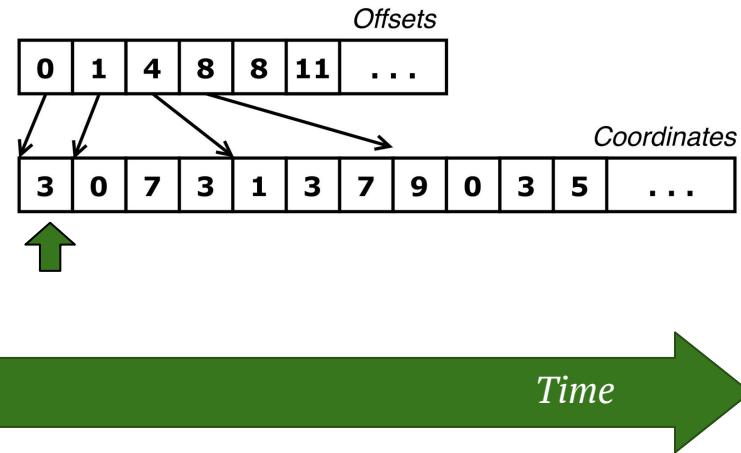
Irregular Accesses Have Poor Temporal And Spatial Locality

```
for v in G:  
    for u in neigh(v):  
        ➔ process(..., vtxData[u], ...)
```

vtxData[3]



LLC

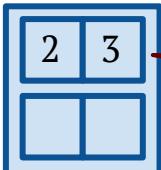


Irregular Accesses Have Poor Temporal And Spatial Locality

```
for v in G:  
    for u in neigh(v):  
        ➔ process(..., vtxData[u], ...)
```

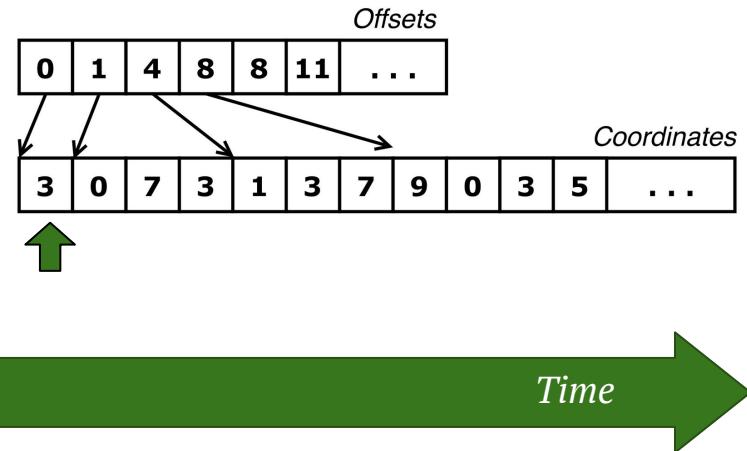
vtxData[3]

Miss



Line granular transfers.
2 words / line

LLC



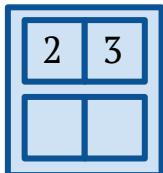
Irregular Accesses Have Poor Temporal And Spatial Locality

```
for v in G:  
    for u in neigh(v):  
        ➔ process(..., vtxData[u], ...)
```

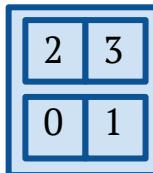
vtxData[3] vtxData[0]



Miss Miss



LLC

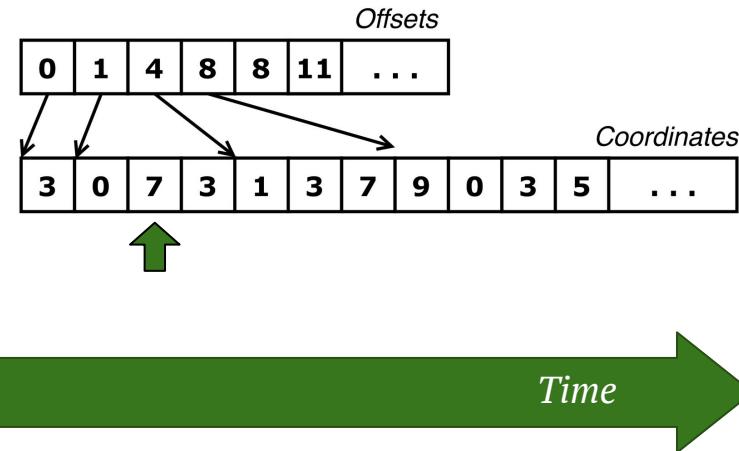
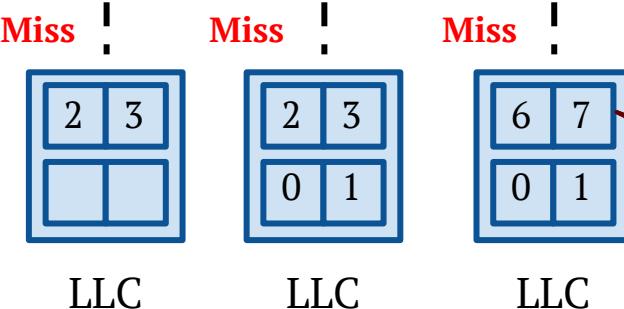


LLC

Irregular Accesses Have Poor Temporal And Spatial Locality

```
for v in G:  
    for u in neigh(v):  
        ➔ process(..., vtxData[u], ...)
```

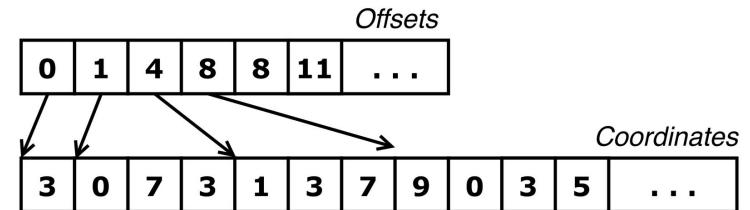
vtxData[3] vtxData[0] vtxData[7]



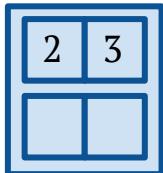
Eviction due to
capacity miss

Irregular Accesses Have Poor Temporal And Spatial Locality

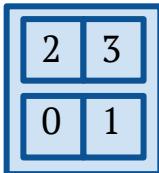
```
for v in G:  
    for u in neigh(v):  
        ➔ process(..., vtxData[u], ...)
```



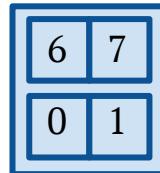
vtxData[3] vtxData[0] vtxData[7] vtxData[3]



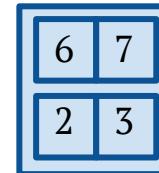
LLC



LLC



LLC

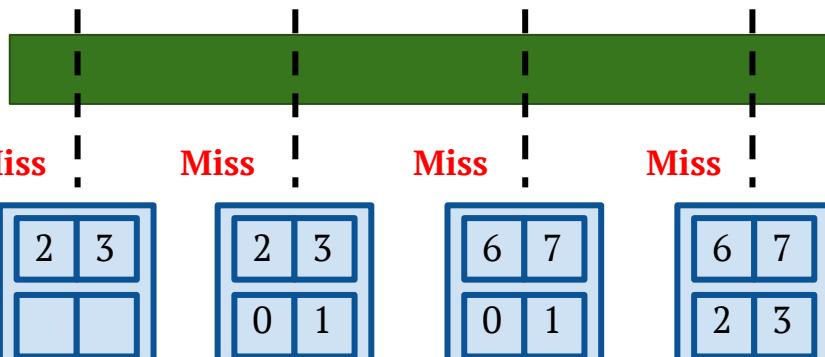


LLC

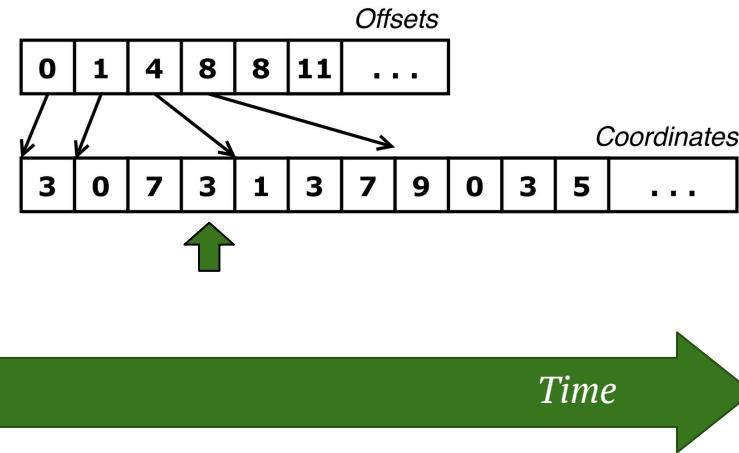
Irregular Accesses Have Poor Temporal And Spatial Locality

```
for v in G:  
    for u in neigh(v):  
        ➔ process(..., vtxData[u], ...)
```

vtxData[3] vtxData[0] vtxData[7] vtxData[3]



LLC

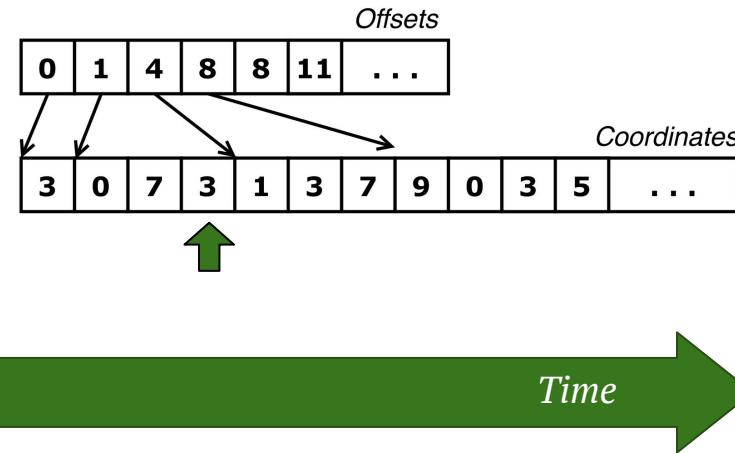
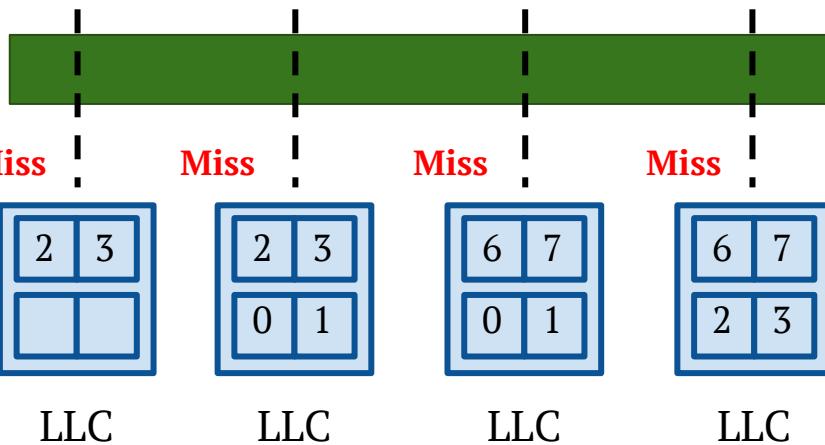


Working set size >> LLC capacity
↓
Poor Temporal Locality

Irregular Accesses Have Poor Temporal And Spatial Locality

```
for v in G:  
    for u in neigh(v):  
        ➔ process(..., vtxData[u], ...)
```

vtxData[3] vtxData[0] vtxData[7] vtxData[3]



Working set size \gg LLC capacity

↓
Poor Temporal Locality

Line Size $>$ Access granularity

↓
Poor Spatial Locality

Outline

- ❖ Poor Locality of Graph Processing Applications ✓
- ❖ Improving locality through Graph Reordering
- ❖ Graph Reordering Challenge - *Application and Input-dependent Speedups*
- ❖ When is Graph Reordering an Optimization?
- ❖ Selective Graph Reordering

Outline

- ❖ Poor Locality of Graph Processing Applications ✓
- ❖ **Improving locality through Graph Reordering**
- ❖ Graph Reordering Challenge - *Application and Input-dependent Speedups*
- ❖ When is Graph Reordering an Optimization?
- ❖ Selective Graph Reordering

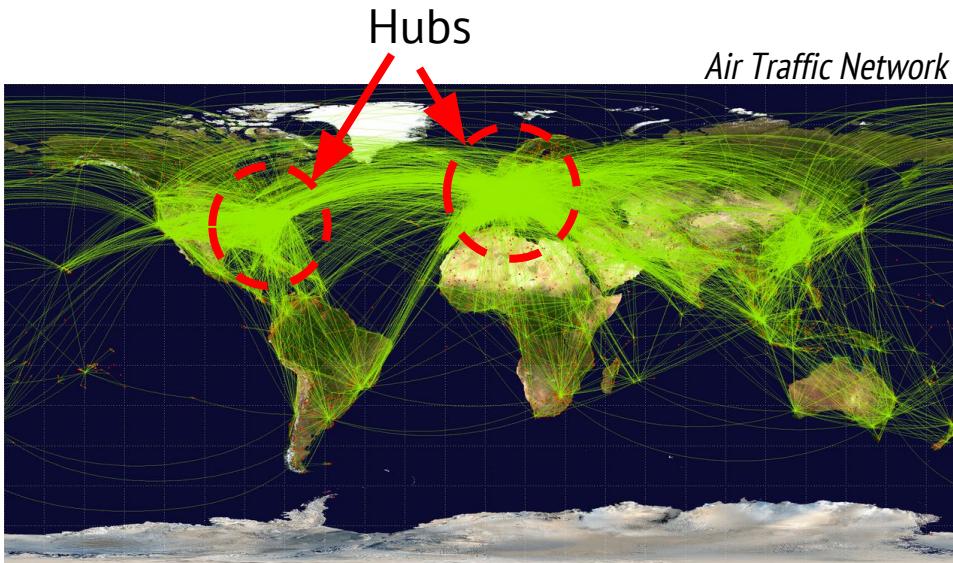
Real-world Graphs Offer Opportunities To Improve Locality

Air Traffic Network



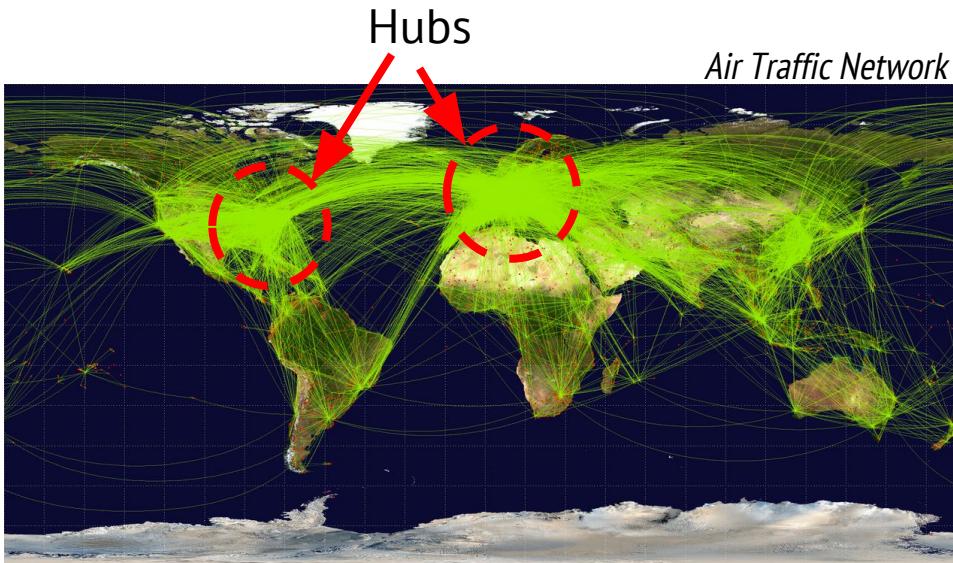
Power-law Degree Distribution

Real-world Graphs Offer Opportunities To Improve Locality



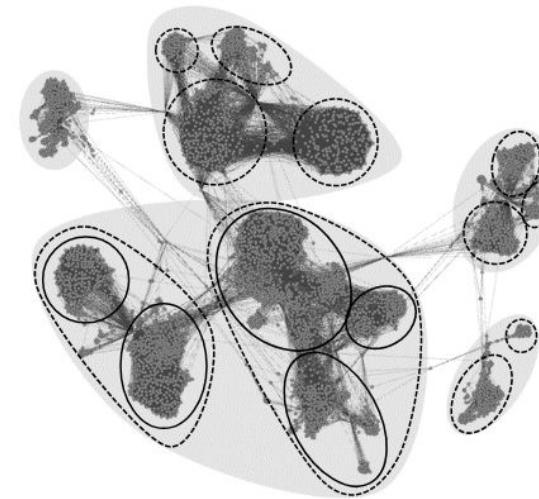
Power-law Degree Distribution

Real-world Graphs Offer Opportunities To Improve Locality



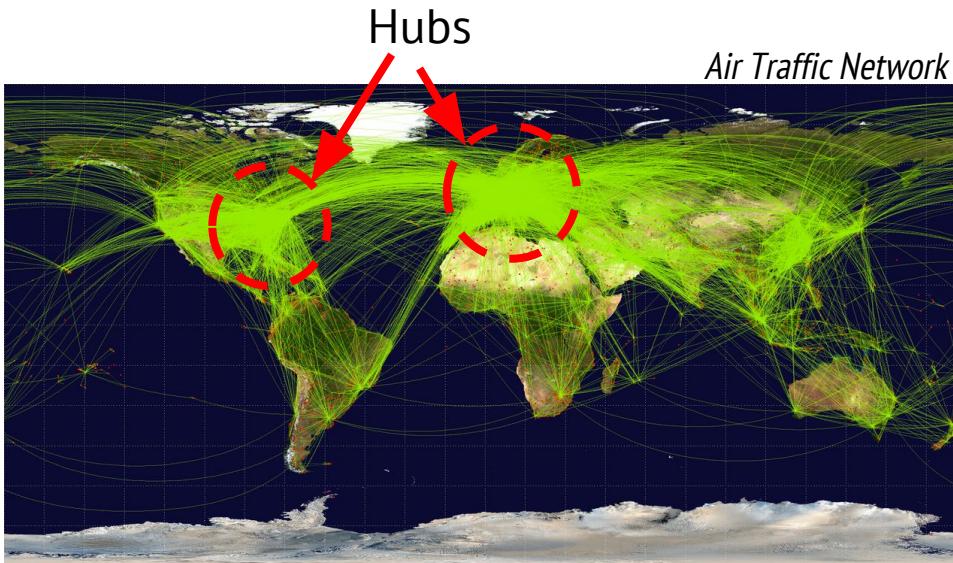
Power-law Degree Distribution

Facebook friend Graph

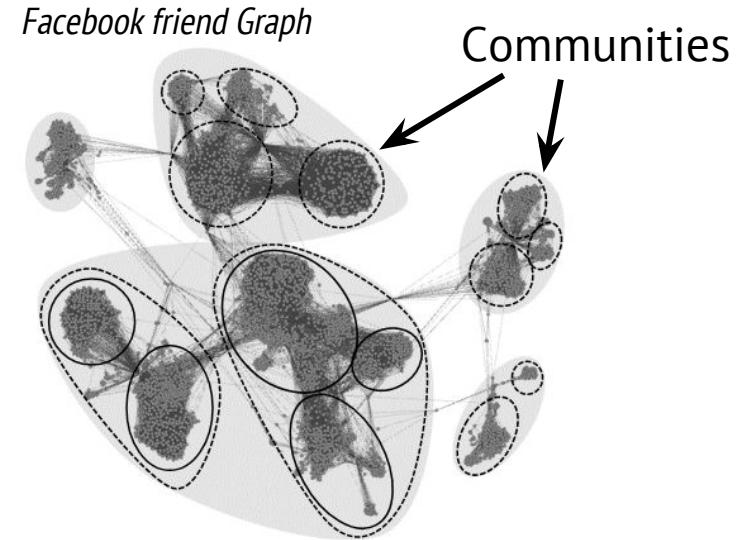


Community Structure

Real-world Graphs Offer Opportunities To Improve Locality

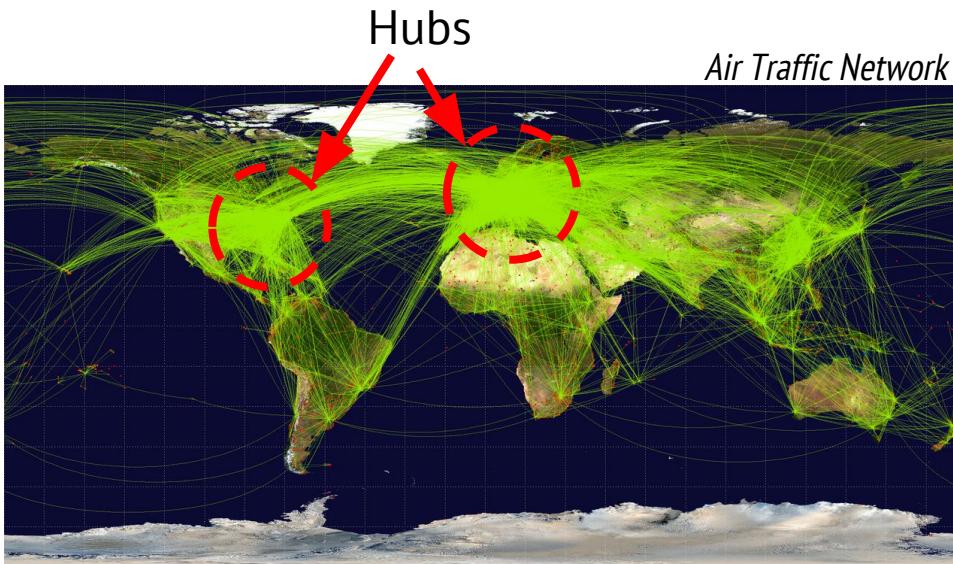


Power-law Degree Distribution

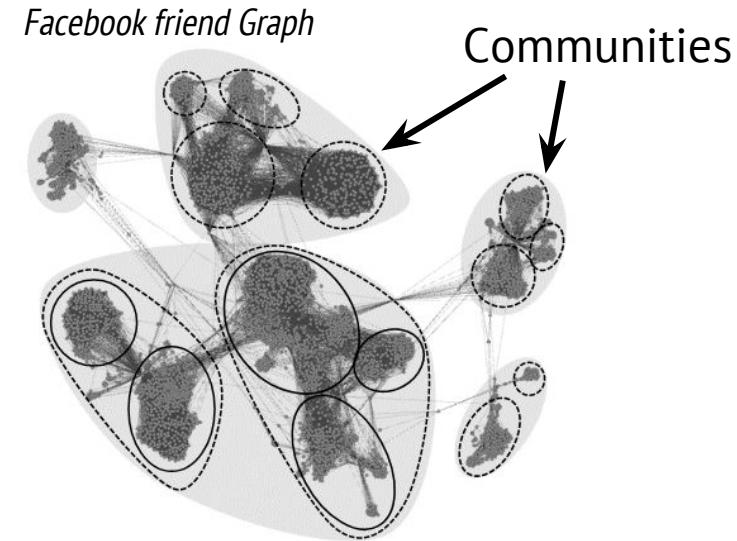


Community Structure

Real-world Graphs Offer Opportunities To Improve Locality



Power-law Degree Distribution



Community Structure

Observation: Subset of vertices are accessed together

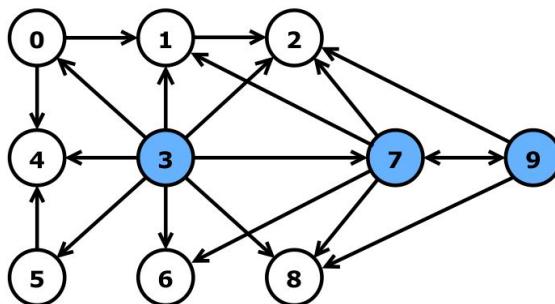
Reordering To Improve Locality of Graph Applications

Key Insight: Store commonly accessed vertices contiguously in memory

Reordering To Improve Locality of Graph Applications

Key Insight: Store commonly accessed vertices contiguously in memory

Power-law graph



Offsets

0	1	4	8	8	11	...
---	---	---	---	---	----	-----

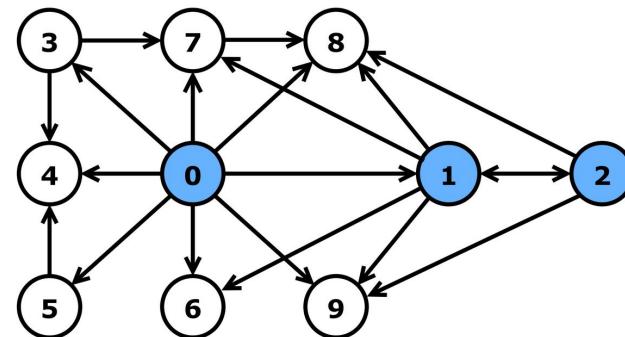
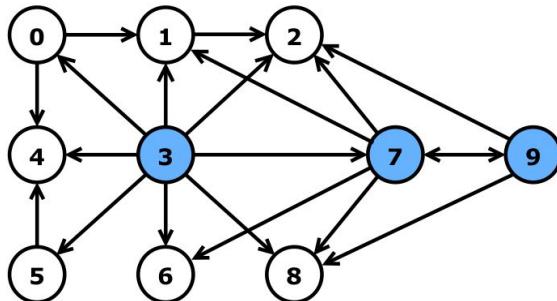
Coordinates

3	0	7	3	1	3	7	9	0	3	5	...
---	---	---	---	---	---	---	---	---	---	---	-----

Reordering To Improve Locality of Graph Applications

Key Insight: Store commonly accessed vertices contiguously in memory

Power-law graph



Offsets

0	1	4	8	8	11	...
---	---	---	---	---	----	-----

Coordinates

3	0	7	3	1	3	7	9	0	3	5	...
---	---	---	---	---	---	---	---	---	---	---	-----

Offsets

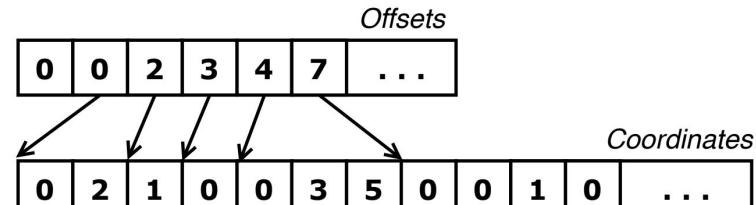
0	0	2	3	4	7	...
---	---	---	---	---	---	-----

Coordinates

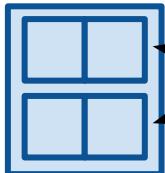
0	2	1	0	0	3	5	0	0	1	0	...
---	---	---	---	---	---	---	---	---	---	---	-----

Reordering Improves Spatial & Temporal Locality

```
for v in G:  
    for u in neigh(v):  
        process(..., vtxData[u], ...)
```



Reordered CSR



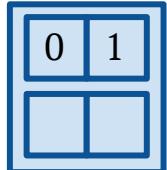
2 lines
2 words/line

LLC

Reordering Improves Spatial & Temporal Locality

```
for v in G:  
    for u in neigh(v):  
        ➔ process(..., vtxData[u], ...)
```

vtxData[0]



LLC

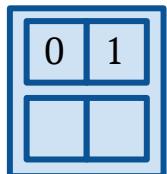
Reordering Improves Spatial & Temporal Locality

```
for v in G:  
    for u in neigh(v):  
        ➔ process(..., vtxData[u], ...)
```

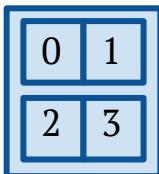
vtxData[0] vtxData[2]



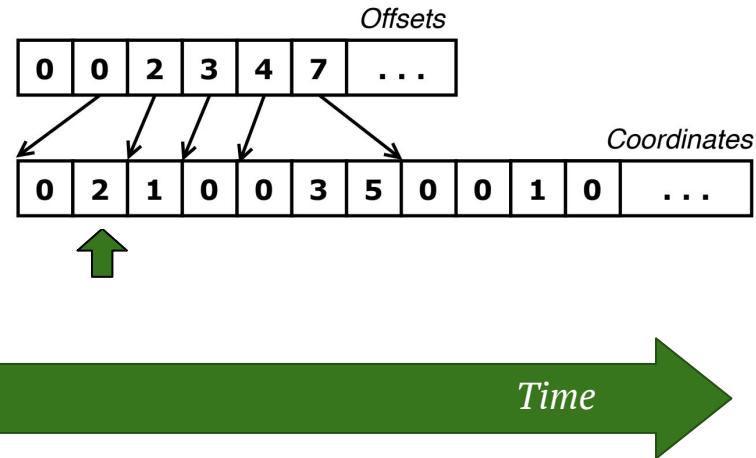
Miss Miss



LLC



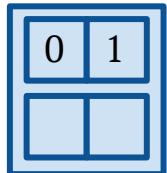
LLC



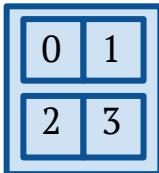
Reordering Improves Spatial & Temporal Locality

```
for v in G:  
    for u in neigh(v):  
        ➔ process(..., vtxData[u], ...)
```

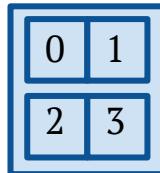
$vtxData[0]$ $vtxData[2]$ $vtxData[1]$



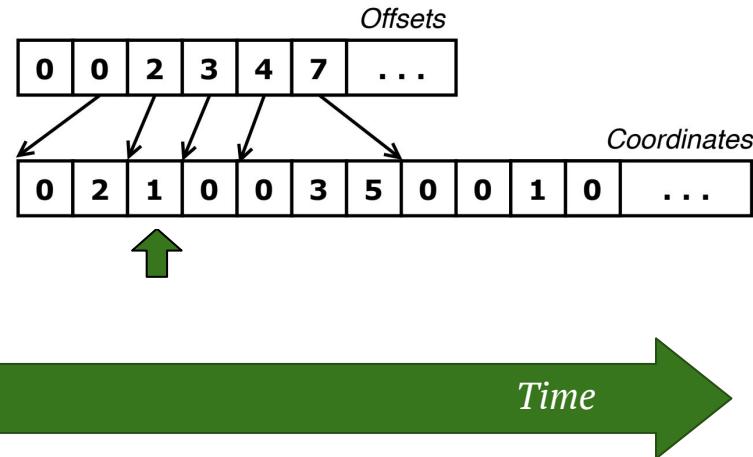
LLC



LLC

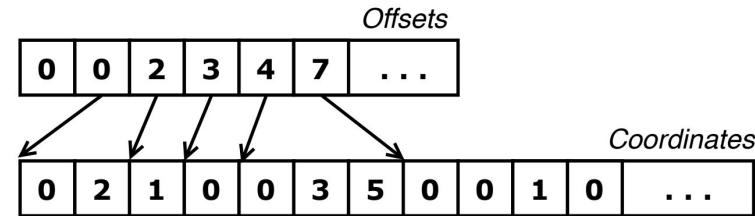


LLC



Reordering Improves Spatial & Temporal Locality

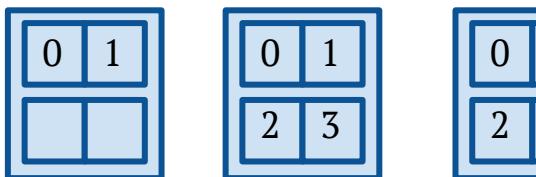
```
for v in G:  
    for u in neigh(v):  
        ➔ process(..., vtxData[u], ...)
```



vtxData[0] *vtxData[2]* *vtxData[1]* *vtxData[0]*



Miss Miss Hit Hit



LLC

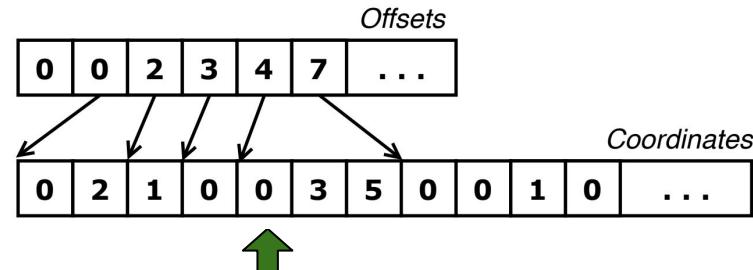
LLC

LLC

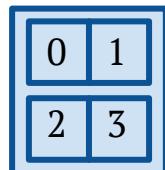
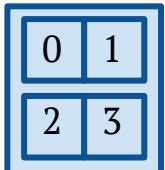
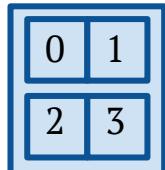
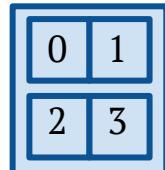
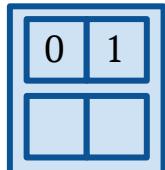
LLC

Reordering Improves Spatial & Temporal Locality

```
for v in G:  
    for u in neigh(v):  
        ➔ process(..., vtxData[u], ...)
```



vtxData[0] vtxData[2] vtxData[1] vtxData[0] vtxData[0]



LLC

LLC

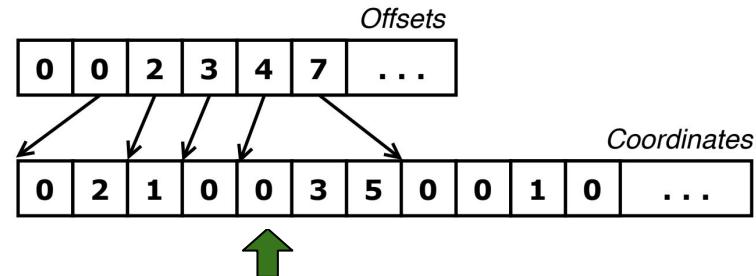
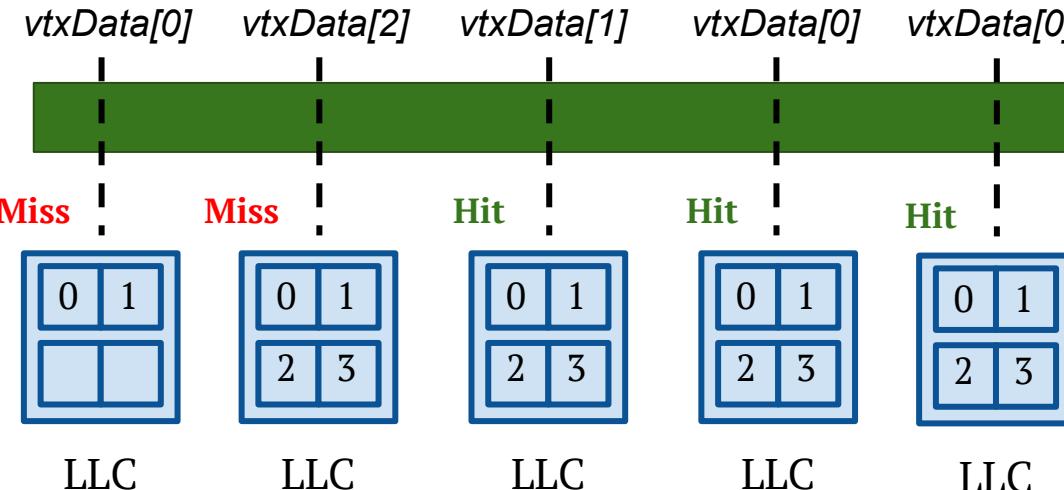
LLC

LLC

LLC

Reordering Improves Spatial & Temporal Locality

```
for v in G:  
    for u in neigh(v):  
        ➔ process(..., vtxData[u], ...)
```



Graph Reordering improved
Spatial and Temporal
locality of `vtxData` accesses

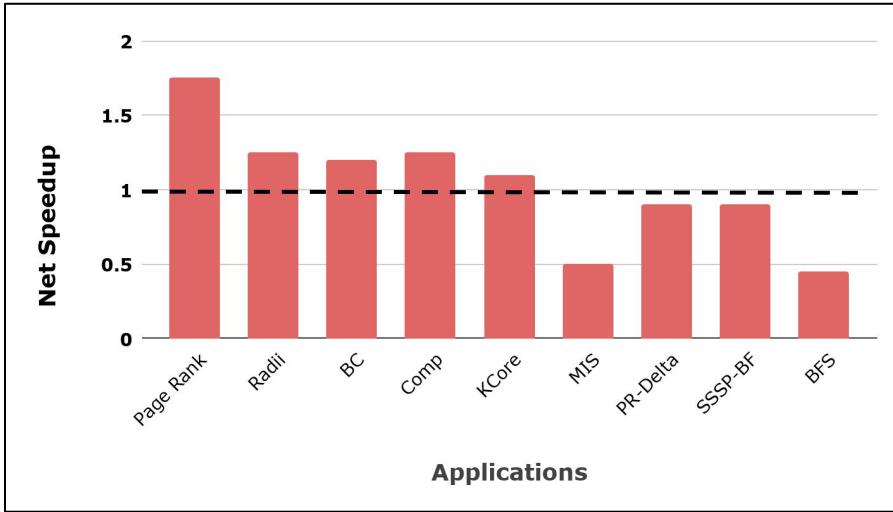
Outline

- ❖ Poor Locality of Graph Processing Applications ✓
- ❖ Improving locality through Graph Reordering ✓
- ❖ Graph Reordering Challenge - *Application and Input-dependent Speedups*
- ❖ When is Graph Reordering an Optimization?
- ❖ Selective Graph Reordering

Outline

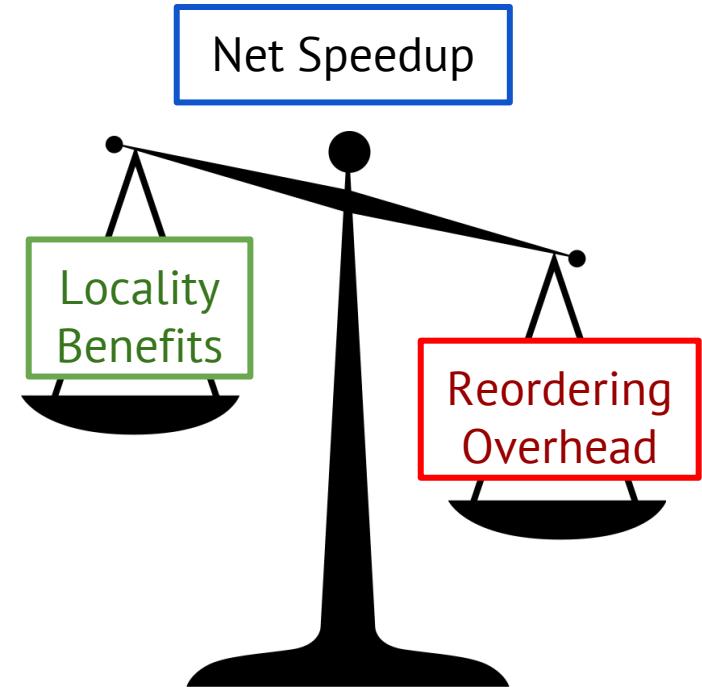
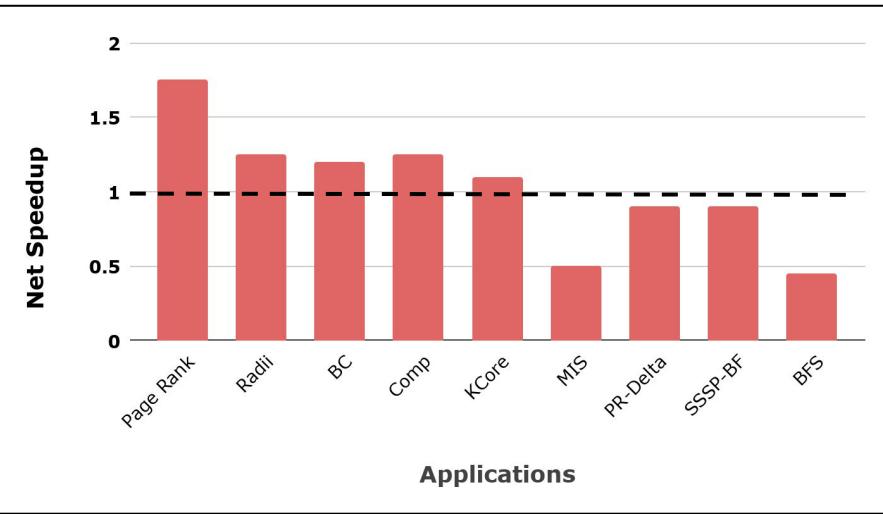
- ❖ Poor Locality of Graph Processing Applications ✓
- ❖ Improving locality through Graph Reordering ✓
- ❖ **Graph Reordering Challenge - *Application and Input-dependent Speedups***
- ❖ When is Graph Reordering an Optimization?
- ❖ Selective Graph Reordering

Graph Reordering Is Not A Panacea



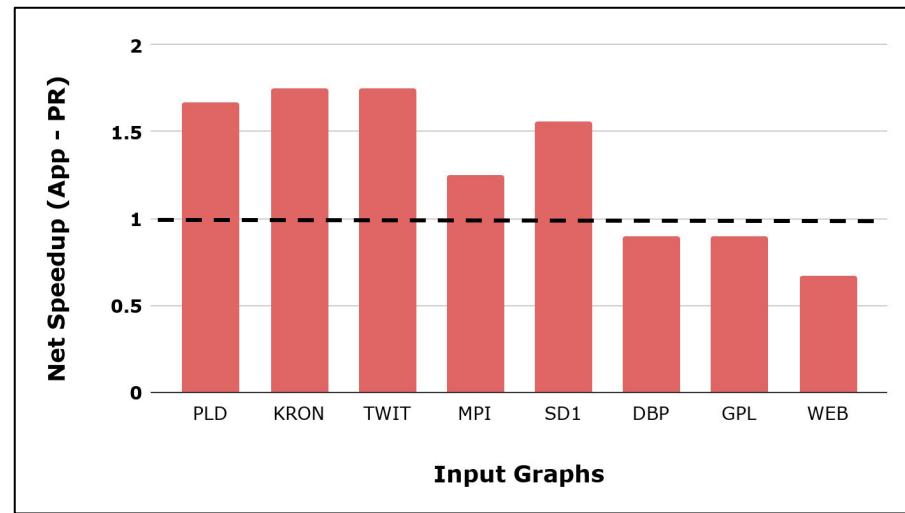
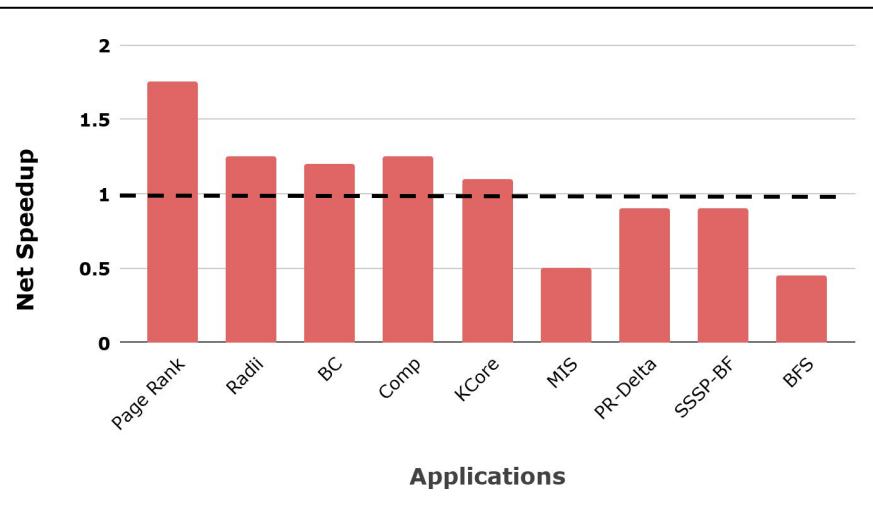
$$\text{Speedup} = \frac{T_{Original}}{T_{Reordered} + \text{Reordering Time}}$$

Graph Reordering Is Not A Panacea



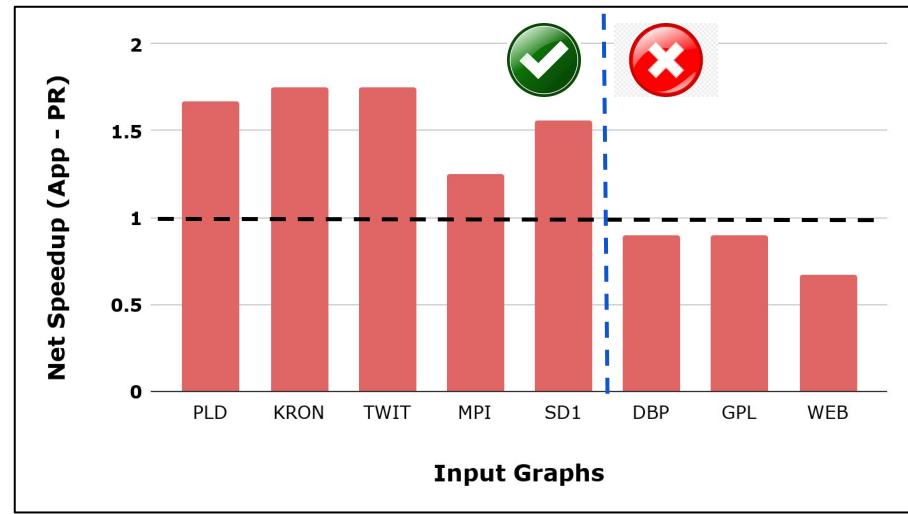
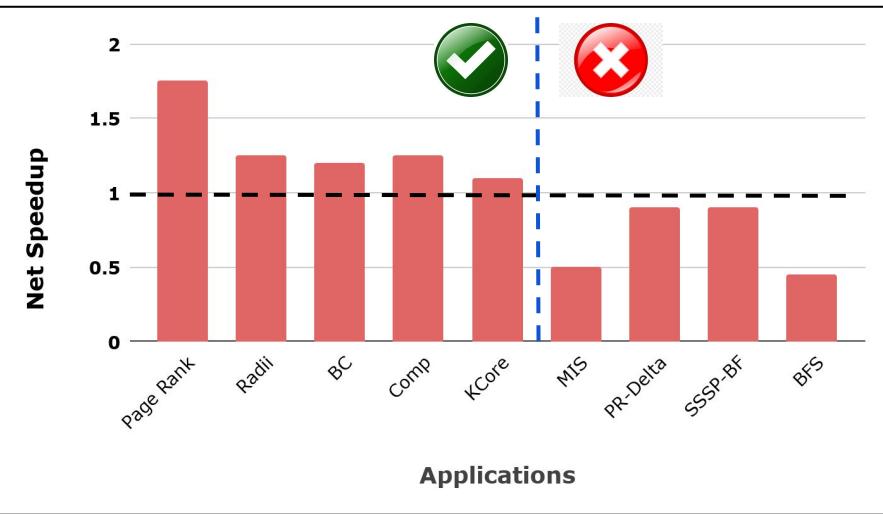
$$\text{Speedup} = \frac{T_{\text{Original}}}{T_{\text{Reordered}} + \text{Reordering Time}}$$

Graph Reordering Is Not A Panacea



$$\text{Speedup} = \frac{T_{\text{Original}}}{T_{\text{Reordered}} + \text{Reordering Time}}$$

Graph Reordering Is Not A Panacea



Net speedup from Reordering depends on the
Application and *Input Graph*

Question: What are the properties of Applications and Input Graphs that benefit from Reordering?

Outline

- ❖ Poor Locality of Graph Processing Applications ✓
- ❖ Improving locality through Graph Reordering ✓
- ❖ Graph Reordering Challenge - *Application and Input-dependent Speedups* ✓
- ❖ When is Graph Reordering an Optimization?
- ❖ Selective Graph Reordering

Outline

- ❖ Poor Locality of Graph Processing Applications ✓
- ❖ Improving locality through Graph Reordering ✓
- ❖ Graph Reordering Challenge - *Application and Input-dependent Speedups* ✓
- ❖ **When is Graph Reordering an Optimization?**
 - Characterization Space
 - Which Applications benefit from Reordering?
 - Which Input Graphs benefit from Reordering?
- ❖ Selective Graph Reordering

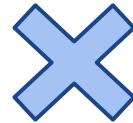
Outline

- ❖ Poor Locality of Graph Processing Applications ✓
- ❖ Improving locality through Graph Reordering ✓
- ❖ Graph Reordering Challenge - *Application and Input-dependent Speedups* ✓
- ❖ **When is Graph Reordering an Optimization?**
 - Characterization Space
 - Which Applications benefit from Reordering?
 - Which Input Graphs benefit from Reordering?
- ❖ Selective Graph Reordering

Characterization Space

3

Graph
Reordering
Techniques



15

Applications
(Ligra, GAP)



8

Input Graphs
(M vertices, B edges)

Server-class Processor

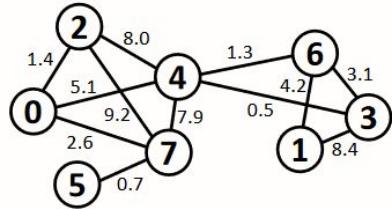
(dual-Socket, 28 cores, 35MB LLC, 64GB DRAM)

Lightweight Reordering (LWR) Techniques

Selection Criteria: Low reordering overhead
(Require very few runs/iterations to amortize overheads)

- Rabbit Ordering [Arai et. al., IPDPS 2016]
- Frequency-based Clustering (or “Hub-Sorting”) [Zhang et. al., Big Data 2017]
- Hub-Clustering (*Our Variation of Hub Sorting*)

LWR 1 - Rabbit Ordering

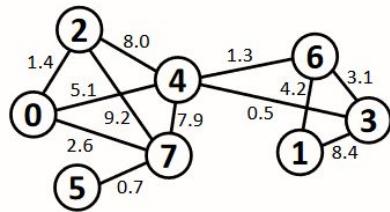


(a) Randomly ordered graph

	0	1	2	3	4	5	6	7
0		1.4	5.1	2.6				
1			8.4		4.2			
2	1.4			8.0		9.2		
3	8.4			0.5	3.1			
4	5.1	8.0	0.5		1.3	7.9		
5							0.7	
6	4.2		3.1	1.3				
7	2.6	9.2	7.9	0.7				

(c) Adjacency matrix of graph (a)

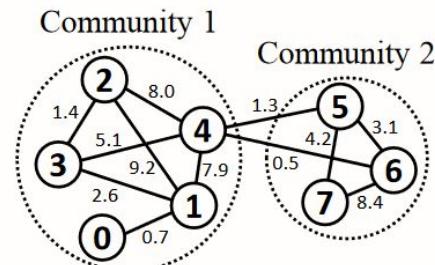
LWR 1 - Rabbit Ordering



(a) Randomly ordered graph

	0	1	2	3	4	5	6	7
0		1.4	5.1		2.6			
1			8.4		4.2			
2	1.4			8.0		9.2		
3		8.4		0.5		3.1		
4	5.1		8.0	0.5		1.3	7.9	
5							0.7	
6		4.2		3.1	1.3			
7	2.6		9.2		7.9	0.7		

(c) Adjacency matrix of graph (a)

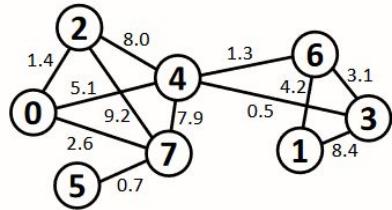


(b) Reordered graph

	0	1	2	3	4	5	6	7
0		0.7						
1	0.7		9.2	2.6	7.9			
2			9.2		1.4	8.0		
3				2.6	1.4		5.1	
4				7.9	8.0	5.1		1.3
5							1.3	3.1
6							0.5	8.4
7							4.2	8.4

(d) Adjacency matrix of graph (b)

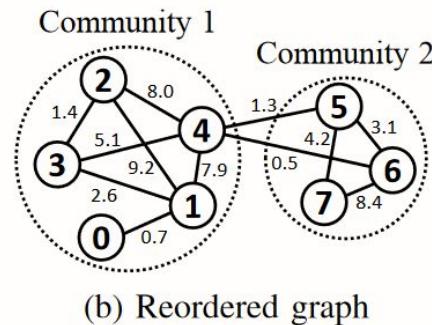
LWR 1 - Rabbit Ordering



(a) Randomly ordered graph

	0	1	2	3	4	5	6	7
0		1.4	5.1		2.6			
1			8.4		4.2			
2	1.4			8.0		9.2		
3		8.4		0.5		3.1		
4	5.1		8.0	0.5		1.3	7.9	
5							0.7	
6		4.2		3.1	1.3			
7	2.6		9.2		7.9	0.7		

(c) Adjacency matrix of graph (a)



(b) Reordered graph

	0	1	2	3	4	5	6	7
0		0.7						
1	0.7		9.2	2.6	7.9			
2		9.2		1.4	8.0			
3			2.6	1.4		5.1		
4			7.9	8.0	5.1		1.3	0.5
5					1.3		3.1	4.2
6					0.5	3.1		8.4
7						4.2	8.4	

(d) Adjacency matrix of graph (b)

- Fast community detection using *incremental aggregation*
- Complexity - $O(|E| - ck|V|)$
where $c = \text{clustering coeff.}$
 $k = \text{avg. degree}$

LWR 2 & 3 - HubSorting & HubClustering



Vertex Degrees (Original)

4	20	4	21	25	99	6	49	64	4
v0	v1	v2	v3	v4	v5	v6	v7	v8	v9

degrees

Vtx IDs

LWR 2 & 3 - HubSorting & HubClustering



Vertex Degrees (Original)

4	20	4	21	25	99	6	49	64	4
v0	v1	v2	v3	v4	v5	v6	v7	v8	v9

degrees

Vtx IDs



Vertex Degrees (Hub Sorted)

99	64	49	21	25	4	6	4	20	4
v0	v1	v2	v3	v4	v5	v6	v7	v8	v9

LWR 2 & 3 - HubSorting & HubClustering



Vertex Degrees (Original)

4	20	4	21	25	99	6	49	64	4
v0	v1	v2	v3	v4	v5	v6	v7	v8	v9

degrees

HubSorting

Vtx IDs

Vertex Degrees (Hub Sorted)

99	64	49	21	25	4	6	4	20	4
v0	v1	v2	v3	v4	v5	v6	v7	v8	v9

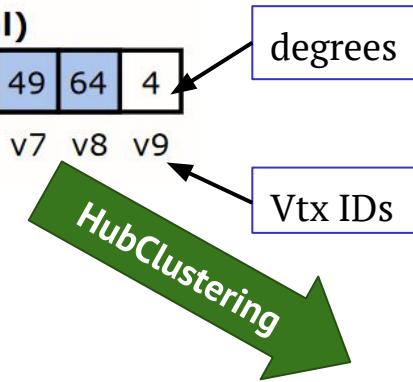
Vertex Degrees (Hub Clustered)

99	49	64	21	25	4	6	20	4	4
v0	v1	v2	v3	v4	v5	v6	v7	v8	v9

LWR 2 & 3 - HubSorting & HubClustering



Vertex Degrees (Original)									
v0	v1	v2	v3	v4	v5	v6	v7	v8	v9
4	20	4	21	25	99	6	49	64	4



Vertex Degrees (Hub Sorted)									
v0	v1	v2	v3	v4	v5	v6	v7	v8	v9
99	64	49	21	25	4	6	4	20	4

Vertex Degrees (Hub Clustered)									
v0	v1	v2	v3	v4	v5	v6	v7	v8	v9
99	49	64	21	25	4	6	20	4	4

	HubSorting	HubClustering
Locality Benefits	Temporal AND Spatial ↑↑	Temporal ↑
Complexity	$O(V \cdot \log V)$ ↓↓	$O(V)$ ↓

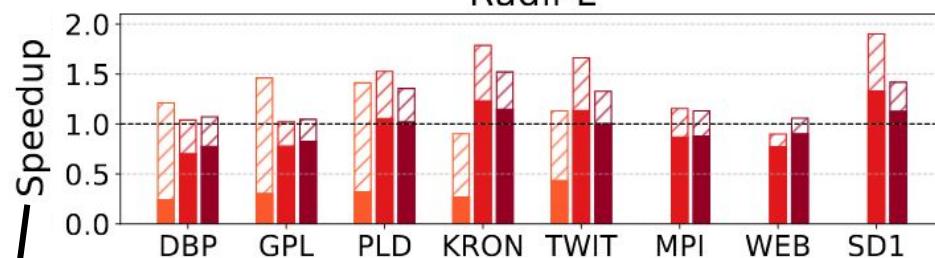
Outline

- ❖ Poor Locality of Graph Processing Applications ✓
- ❖ Improving locality through Graph Reordering ✓
- ❖ Graph Reordering Challenge - *Application and Input-dependent Speedups* ✓
- ❖ **When is Graph Reordering an Optimization?**
 - Characterization Space ✓
 - Which Applications benefit from Reordering?
 - Which Input Graphs benefit from Reordering?
- ❖ Selective Graph Reordering

Outline

- ❖ Poor Locality of Graph Processing Applications ✓
- ❖ Improving locality through Graph Reordering ✓
- ❖ Graph Reordering Challenge - *Application and Input-dependent Speedups* ✓
- ❖ **When is Graph Reordering an Optimization?**
 - Characterization Space ✓
 - Which Applications benefit from Reordering?
 - Which Input Graphs benefit from Reordering?
- ❖ Selective Graph Reordering

Legend for Results



Application

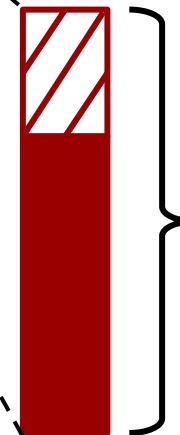
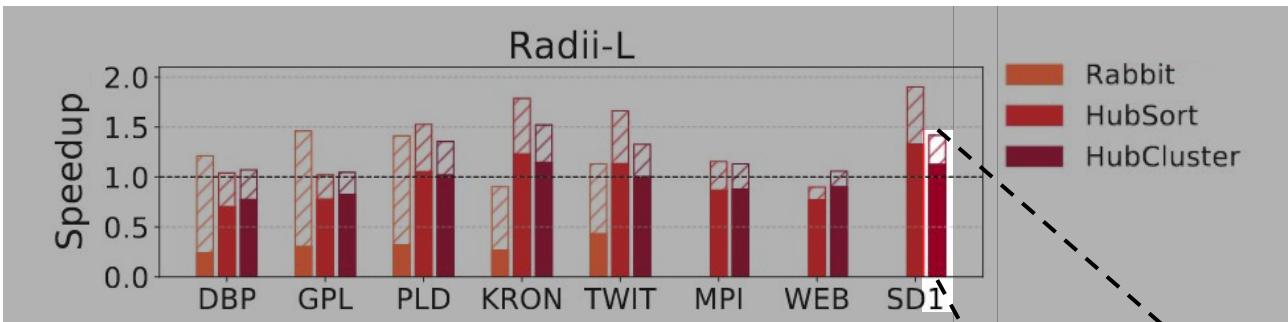
Rabbit
HubSort
HubCluster

Reordering
Technique

Input Graph

Speedup with respect to original ordering of graph

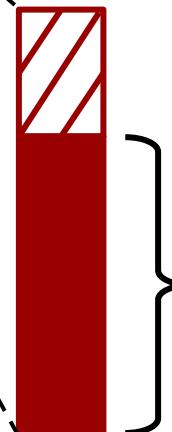
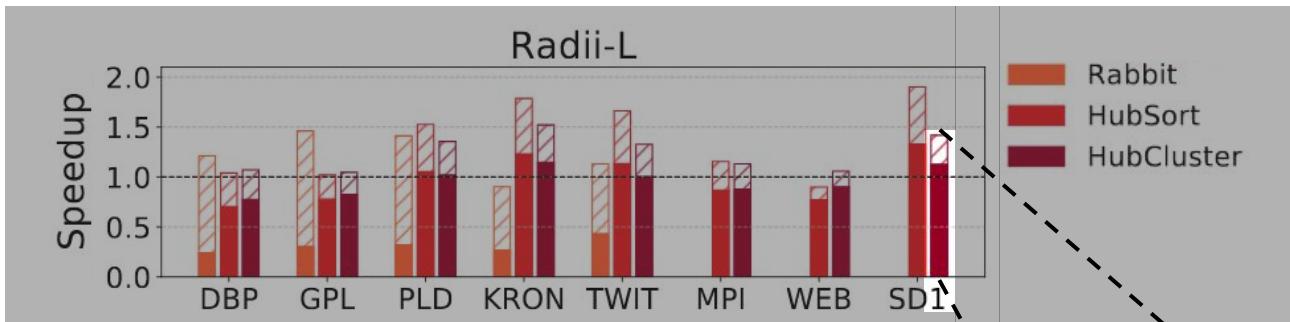
Legend for Results



Speedup excluding
the overhead of
reordering

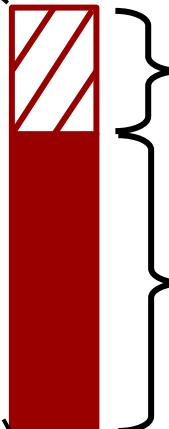
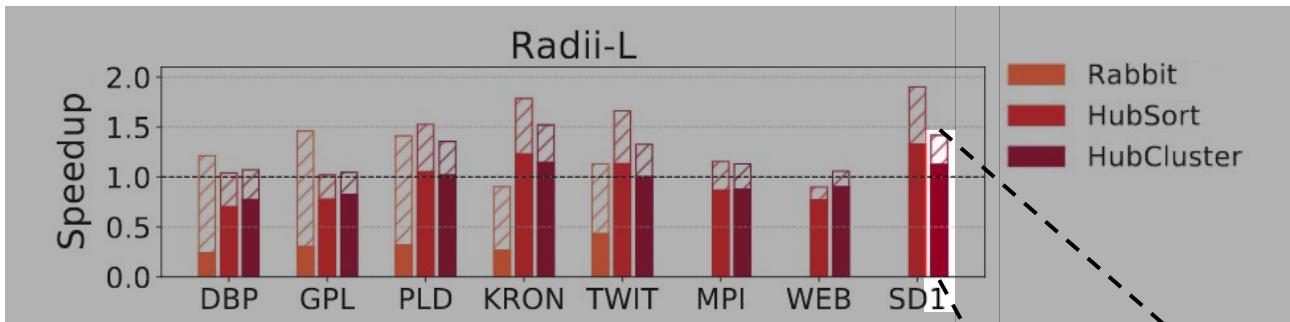
(Max Speedup)

Legend for Results



Speedup including
the overhead of
reordering
(Net Speedup)

Legend for Results



Reordering overhead

Speedup including
the overhead of
reordering

(Net Speedup)

15 Applications → 5 Categories

- **Category 1:** Applications processing Large Frontiers are *good candidates*
- **Category 2:** Symmetric bipartite graphs require *bi-partiteness aware reordering*
- **Category 3:** Applications processing small frontiers *offer limited opportunity*
- **Category 4:** Reordering for Push-style applications introduces *false-sharing*
- **Category 5:** Reordering *affects convergence* for applications with ID-dependent computations

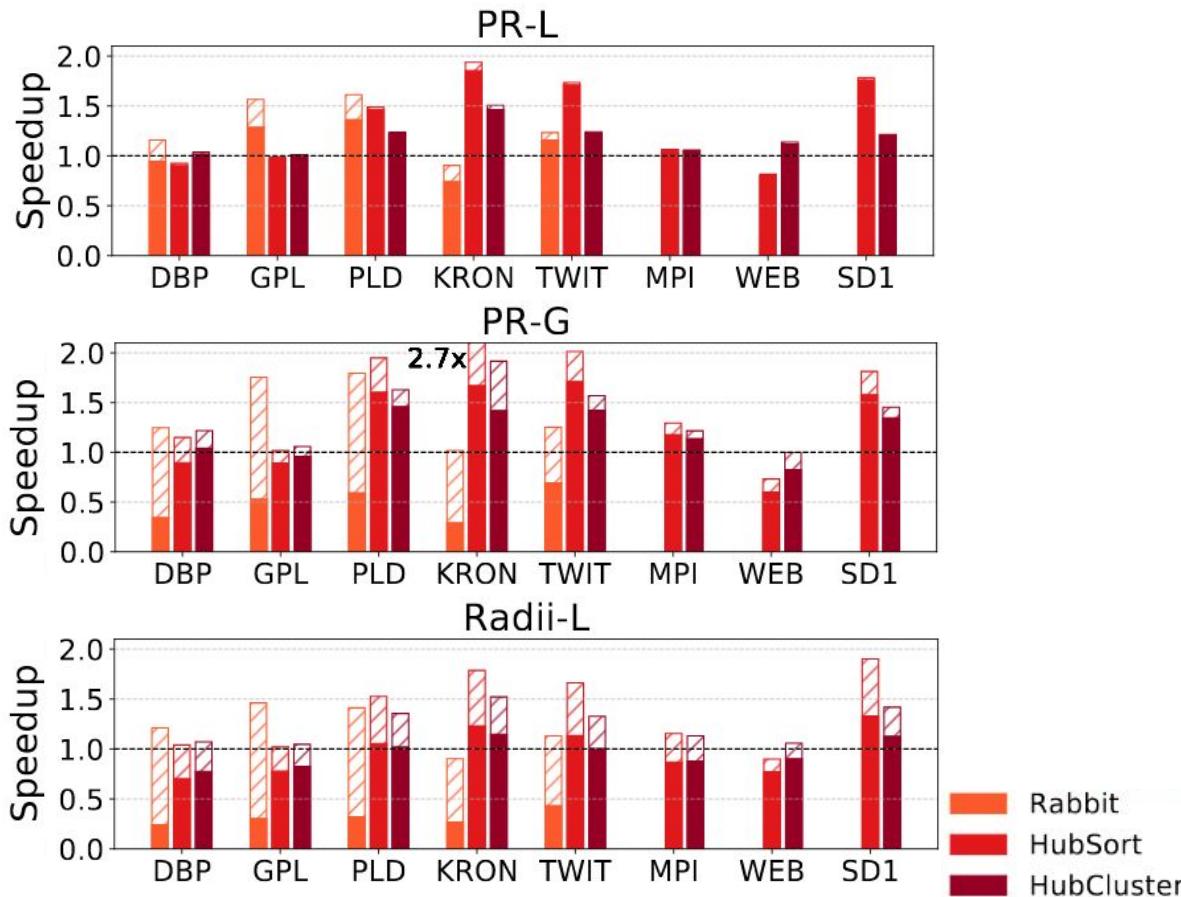
15 Applications → 5 Categories

- **Category 1:** Applications processing Large Frontiers are *good candidates*
- **Category 2:** Symmetric bipartite graphs require *bi-partiteness aware reordering*
- **Category 3:** Applications processing small frontiers *offer limited opportunity*
- **Category 4:** Reordering for Push-style applications introduces *false-sharing*
- **Category 5:** Reordering *affects convergence* for applications with ID-dependent computations

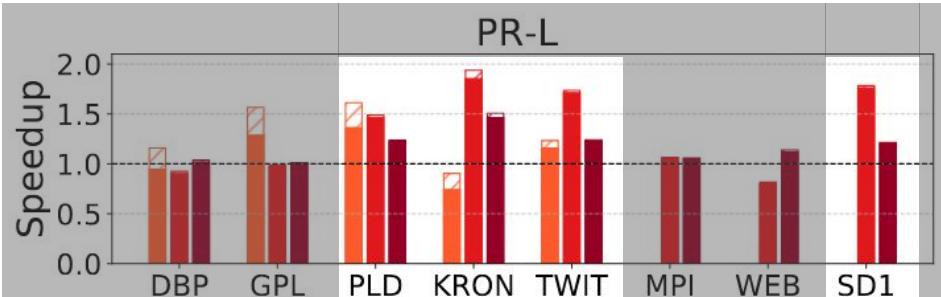
Category I - Applications Processing a Large Fraction Of Edges

- ❖ PageRank (Ligra & Gap)
- ❖ Graph Radii Estimation (Ligra)

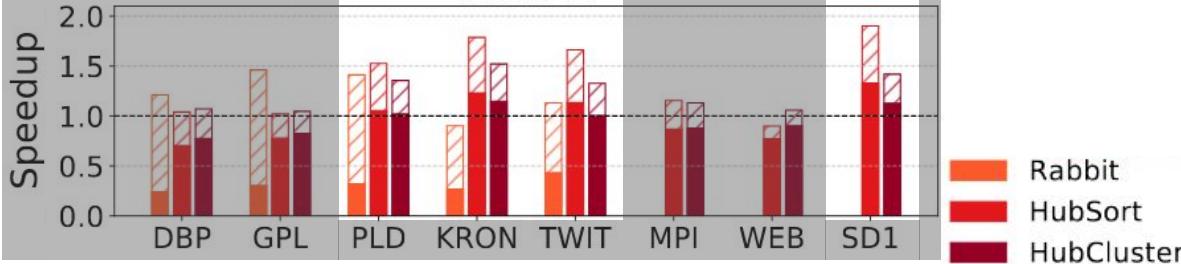
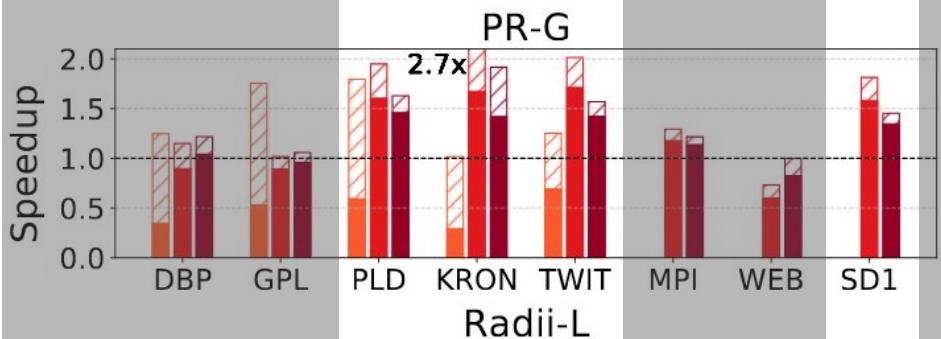
Category I - Applications Processing a Large Fraction Of Edges



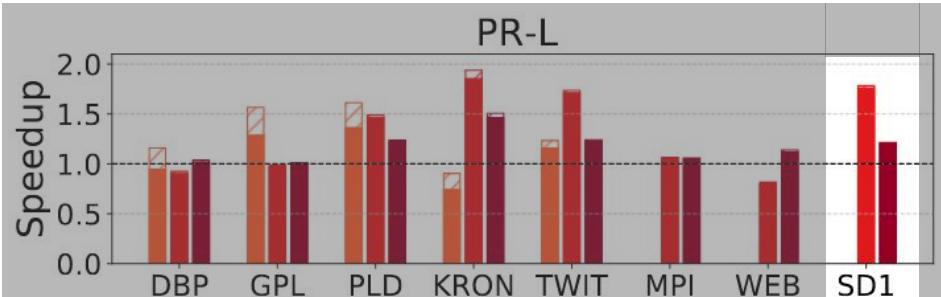
Category I - Applications Processing a Large Fraction Of Edges



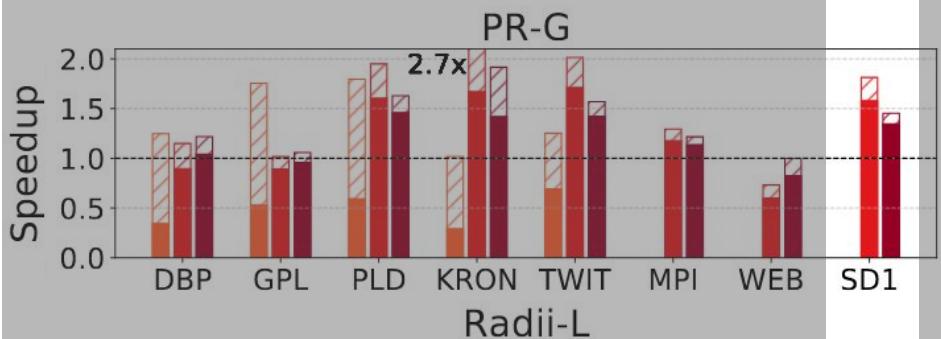
Observation 1: LWR provides *end-to-end* speedups in some cases



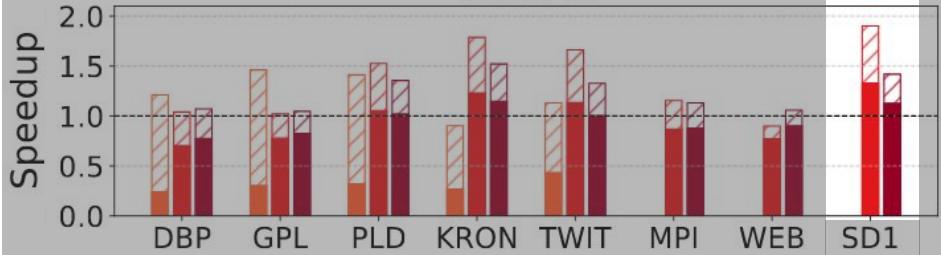
Category I - Applications Processing a Large Fraction Of Edges



Observation 1: LWR provides *end-to-end* speedups in some cases

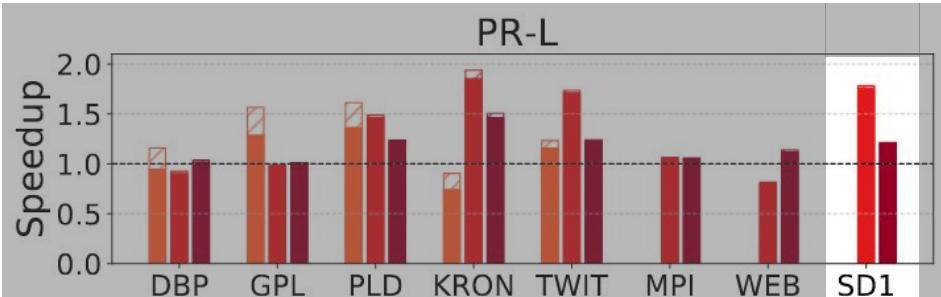


Observation 2: Maximum speedups from HubSort > HubCluster

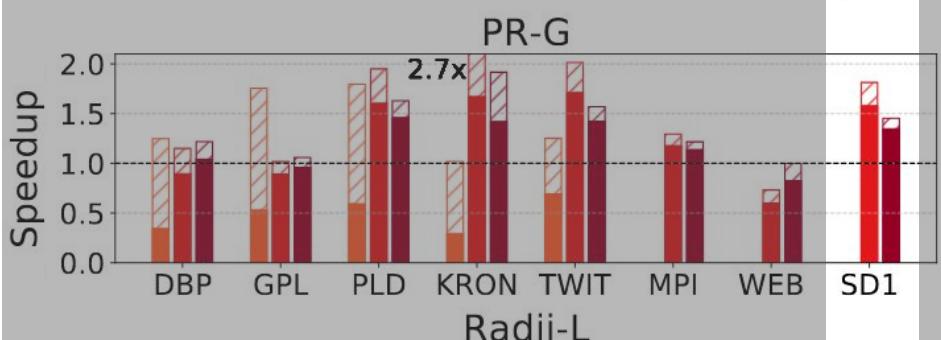


- Rabbit
- HubSort
- HubCluster

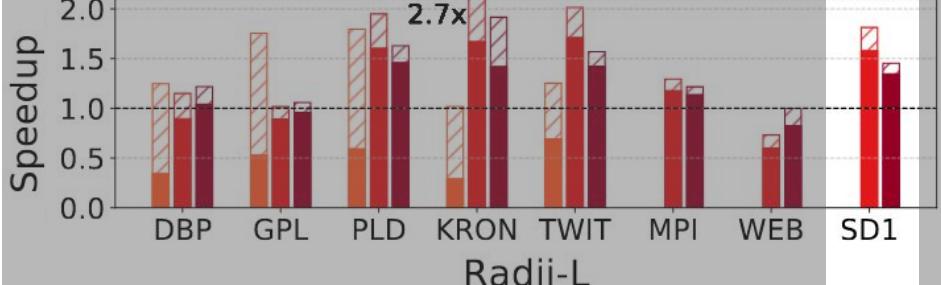
Category I - Applications Processing a Large Fraction Of Edges



Observation 1: LWR provides *end-to-end* speedups in some cases

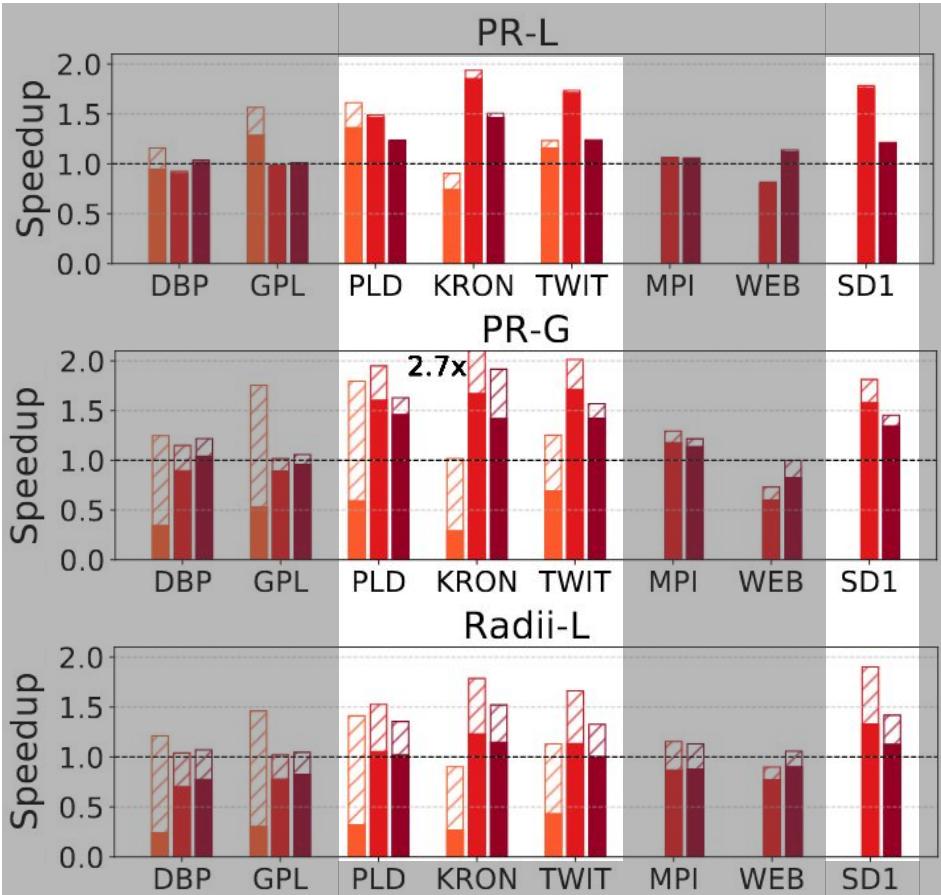


Observation 2: Maximum speedups from HubSort > HubCluster



Observation 3: Reordering Overhead is HubSort > HubCluster

Category I - Applications Processing a Large Fraction Of Edges



Observation 1: LWR provides *end-to-end* speedups in some cases

Observation 2: Maximum speedups from HubSort > HubCluster

Observation 3: Reordering Overhead is HubSort > HubCluster

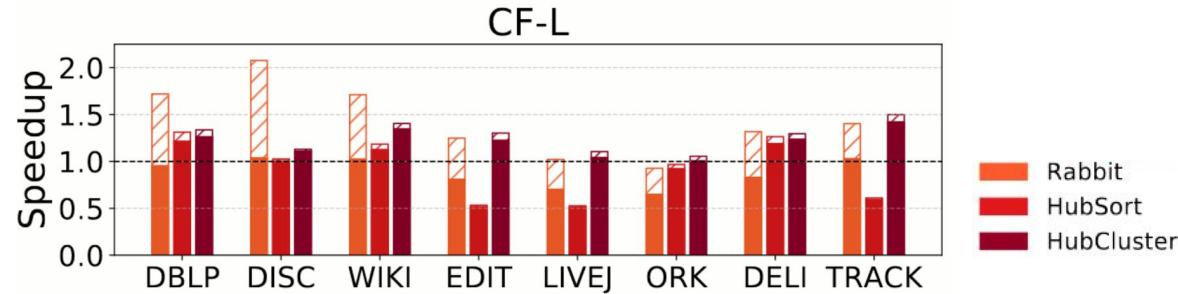
Observation 4: HubSort strikes a balance between effectiveness and overhead

- Rabbit
- HubSort
- HubCluster

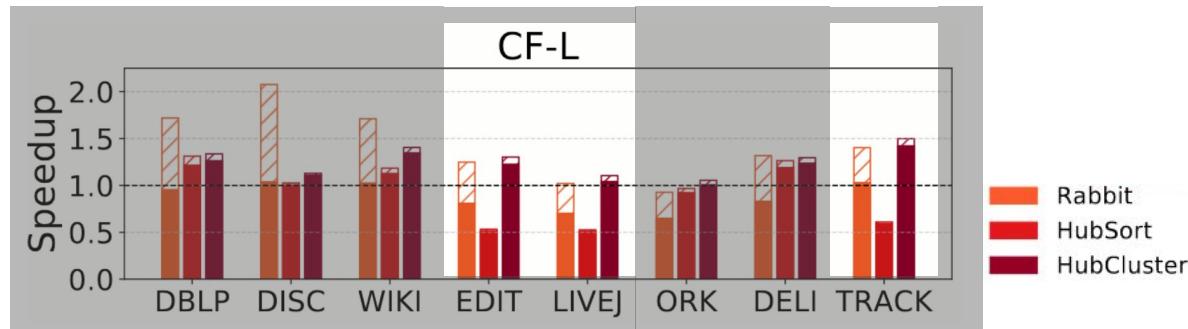
Category II - Executions On Symmetric Bipartite Graphs

- ❖ Collaborative Filtering (Ligra)

Category II - Executions On Symmetric Bipartite Graphs

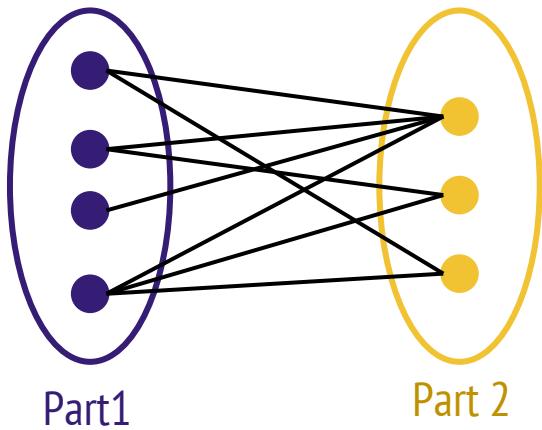


Category II - Executions On Symmetric Bipartite Graphs

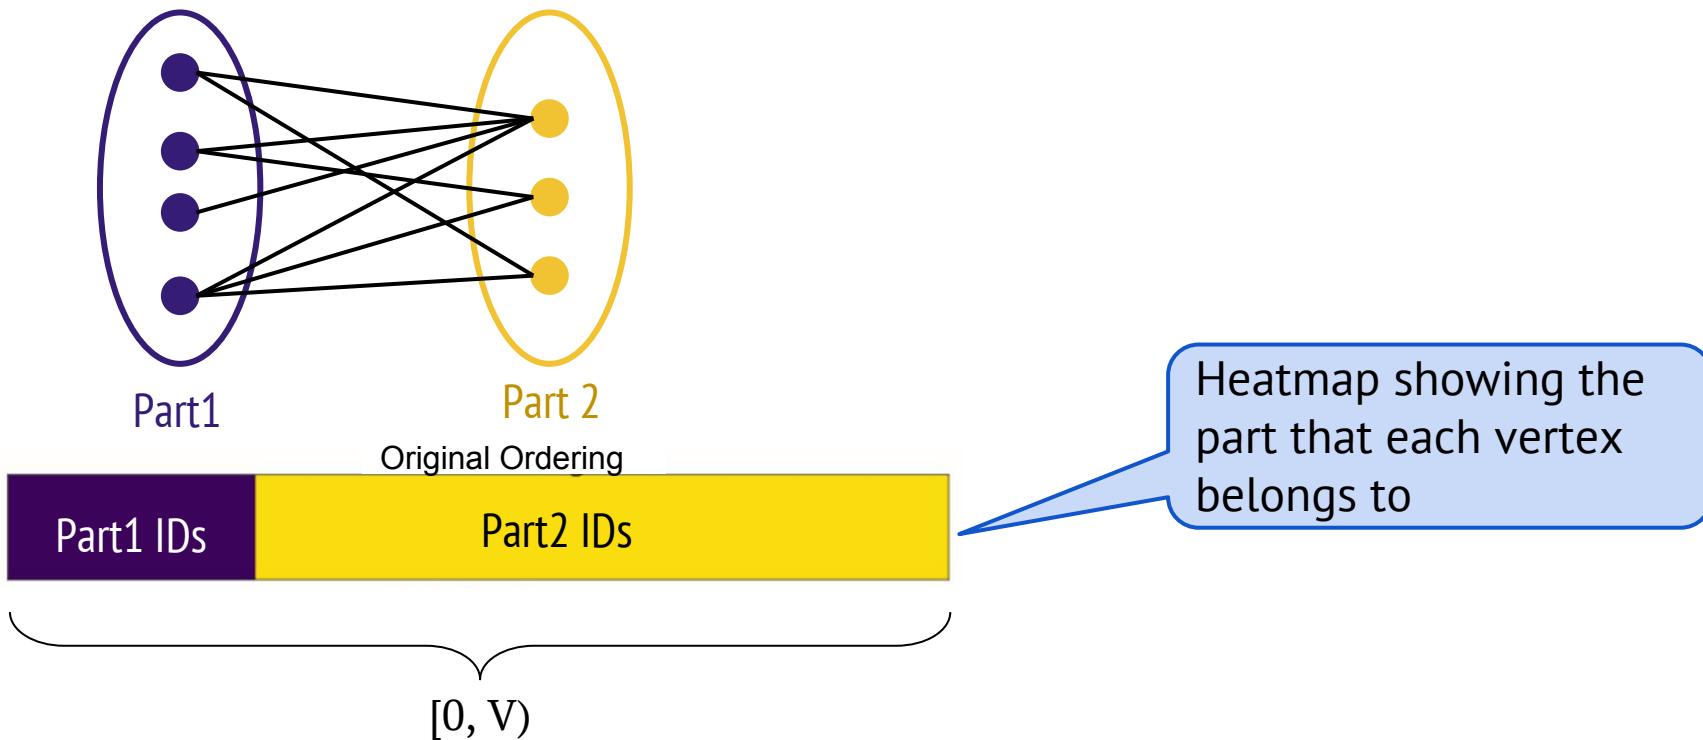


Surprising trend: HubSort causes net slowdowns

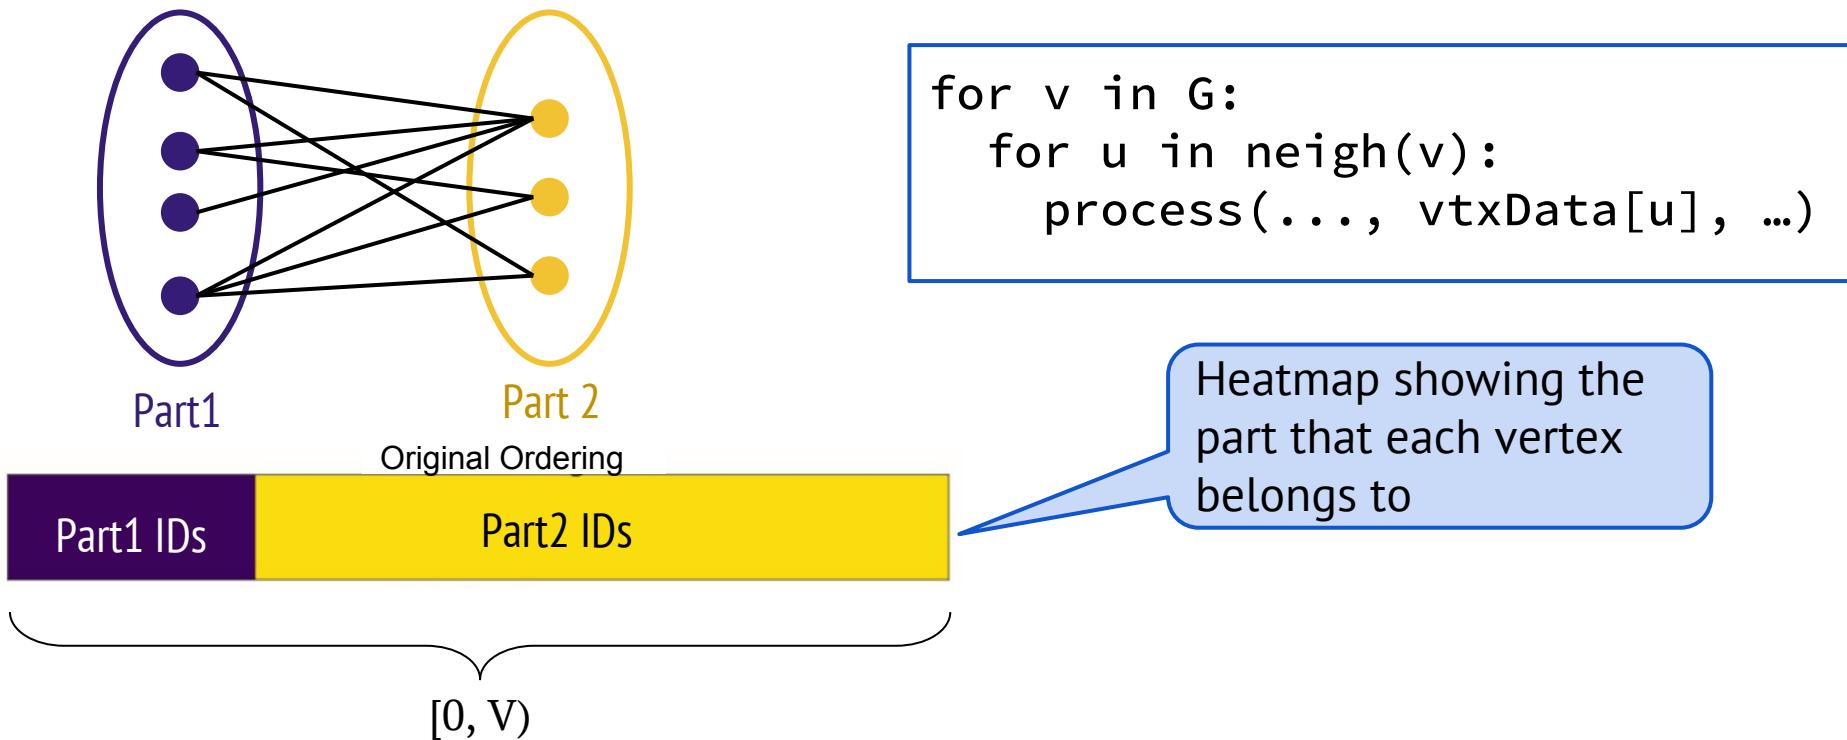
Category II - Reason For Slowdown With HubSort



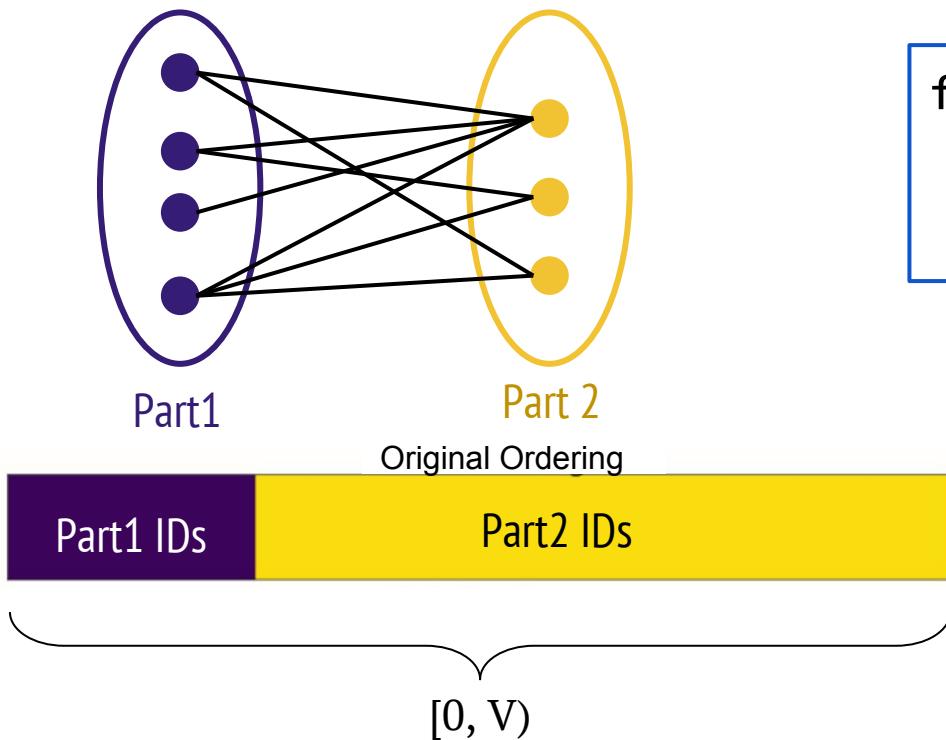
Category II - Reason For Slowdown With HubSort



Category II - Reason For Slowdown With HubSort



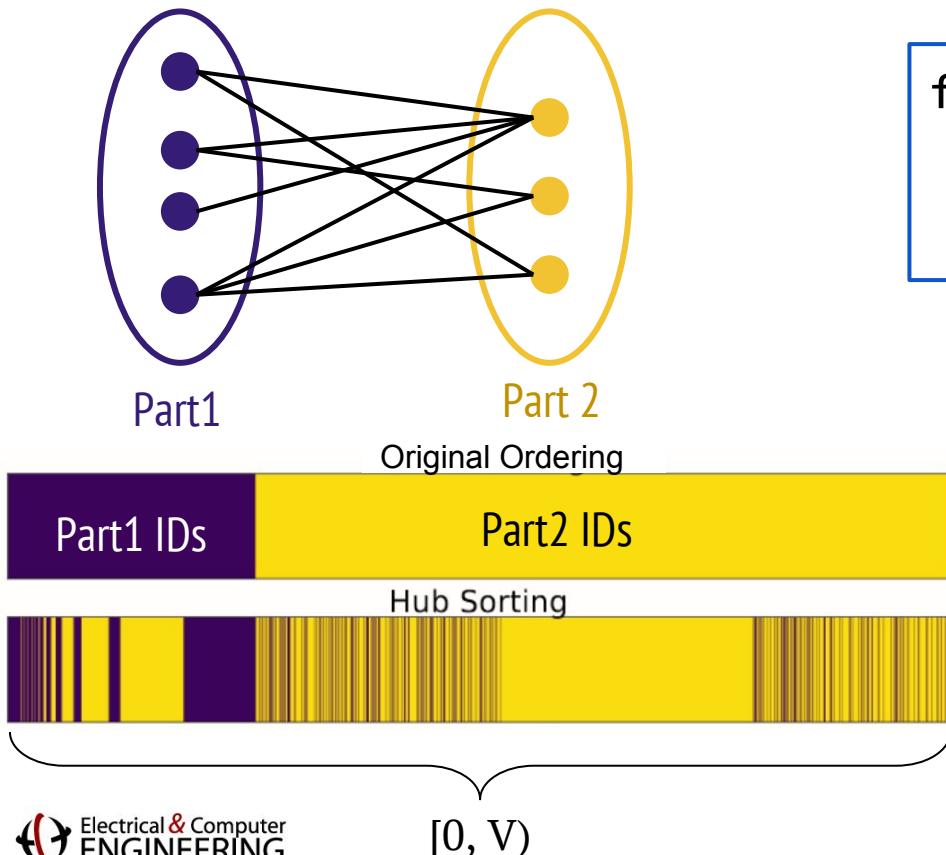
Category II - Reason For Slowdown With HubSort



```
for v in G:  
    for u in neigh(v):  
        process(..., vtxData[u], ...)
```

Range of Irregular accesses
to `vtxData` = #nodes in a part

Category II - Reason For Slowdown With HubSort

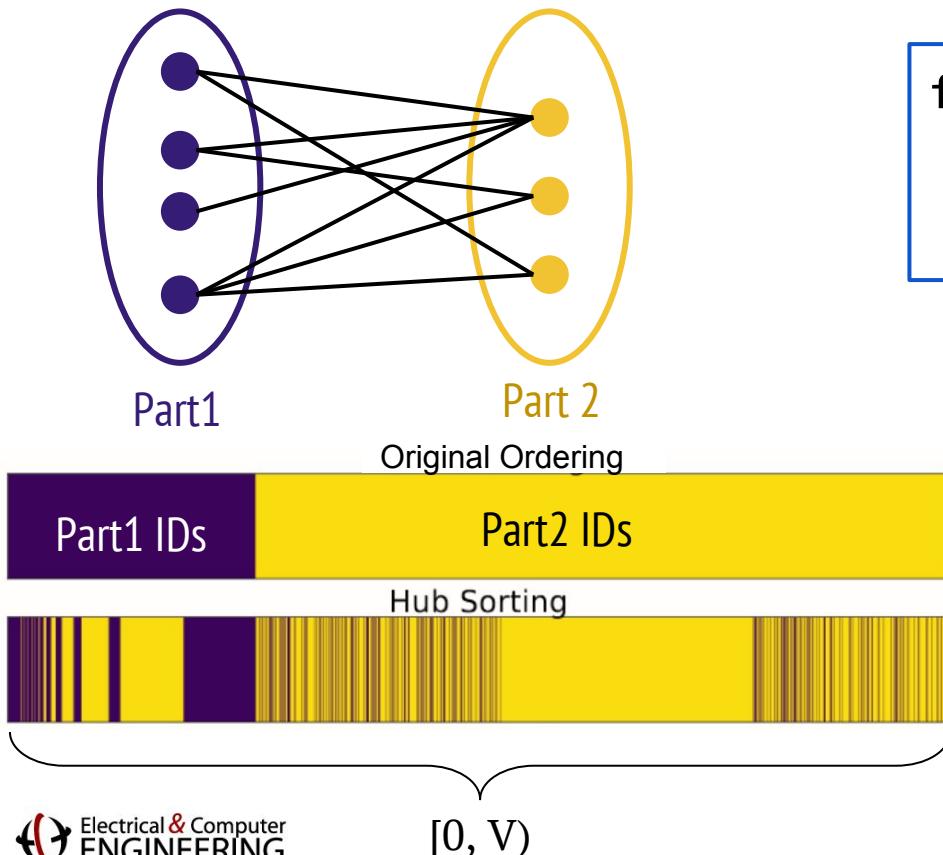


```
for v in G:  
    for u in neigh(v):  
        process(..., vtxData[u], ...)
```

Vertices from different parts assigned consecutive IDs

↓
Increased range of irregular accesses to vtxData

Category II - Reason For Slowdown With HubSort



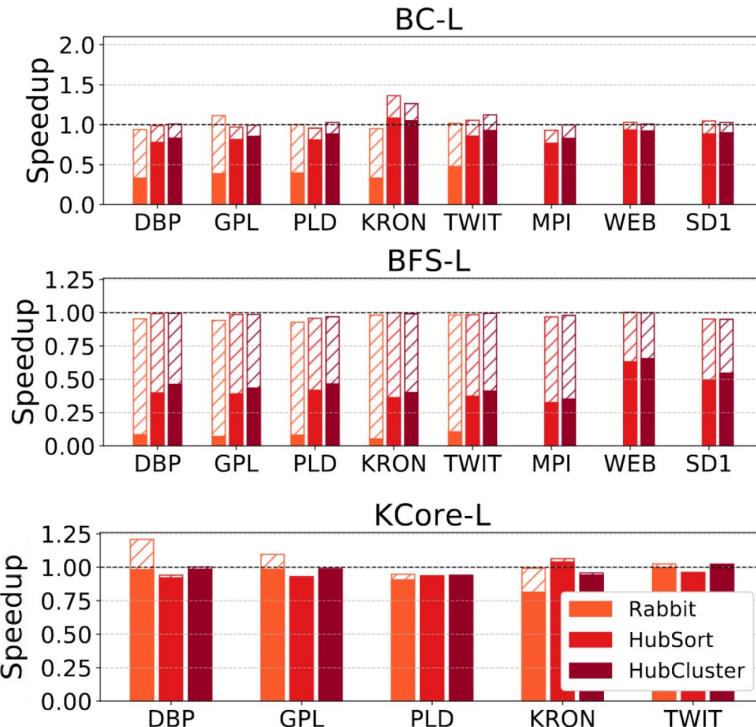
```
for v in G:  
    for u in neigh(v):  
        process(..., vtxData[u], ...)
```

Opportunity: Bipartiteness-aware
Graph Reordering Techniques

Category III - Applications Processing a Small Fraction of Edges

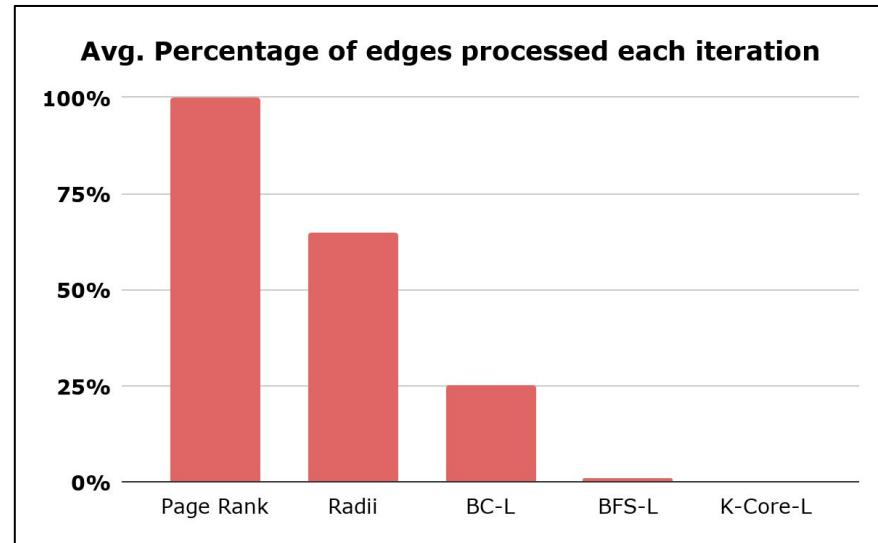
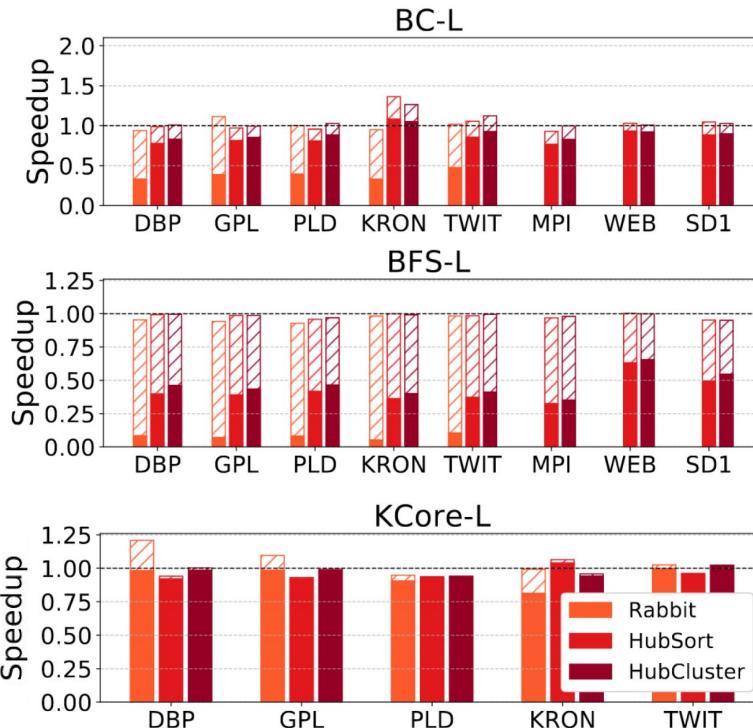
- ❖ Betweenness Centrality
- ❖ BFS
- ❖ K-Core Decomposition

Category III - Applications Processing a Small Fraction of Edges

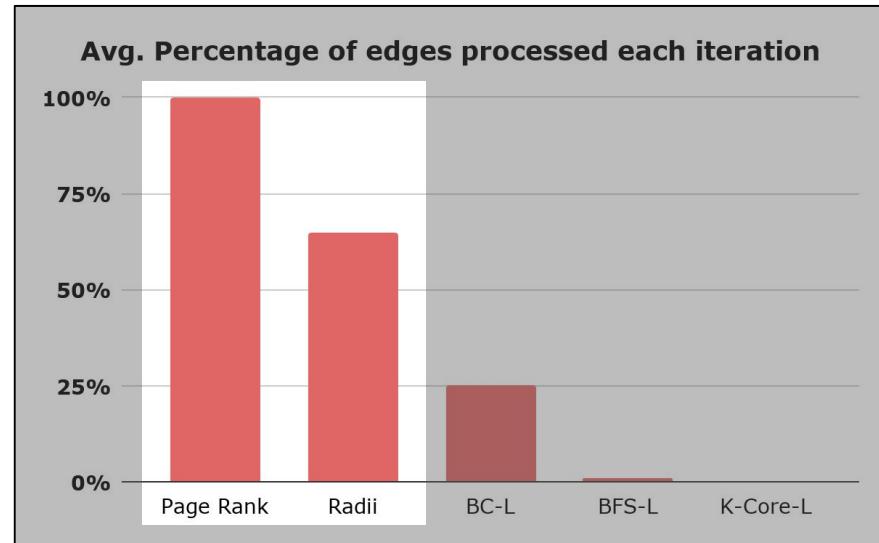
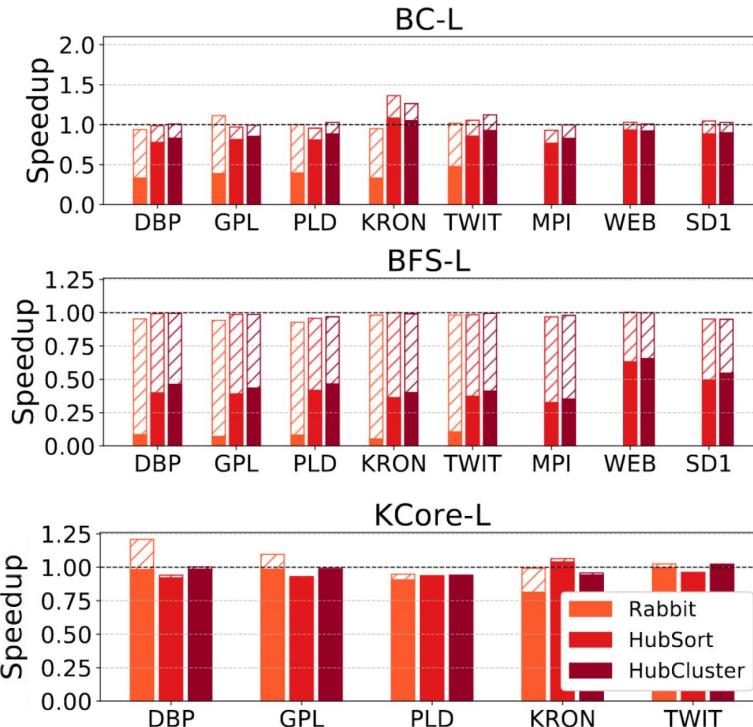


Low speedup even without
Reordering overheads

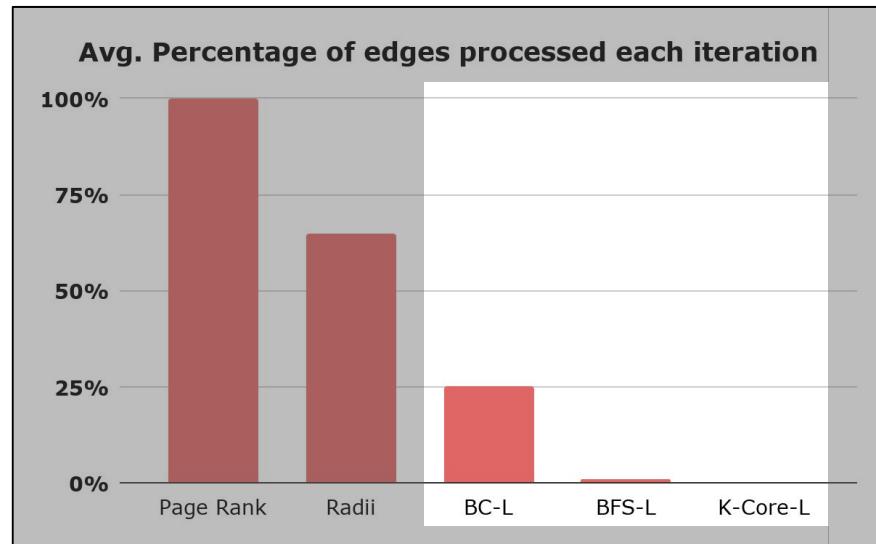
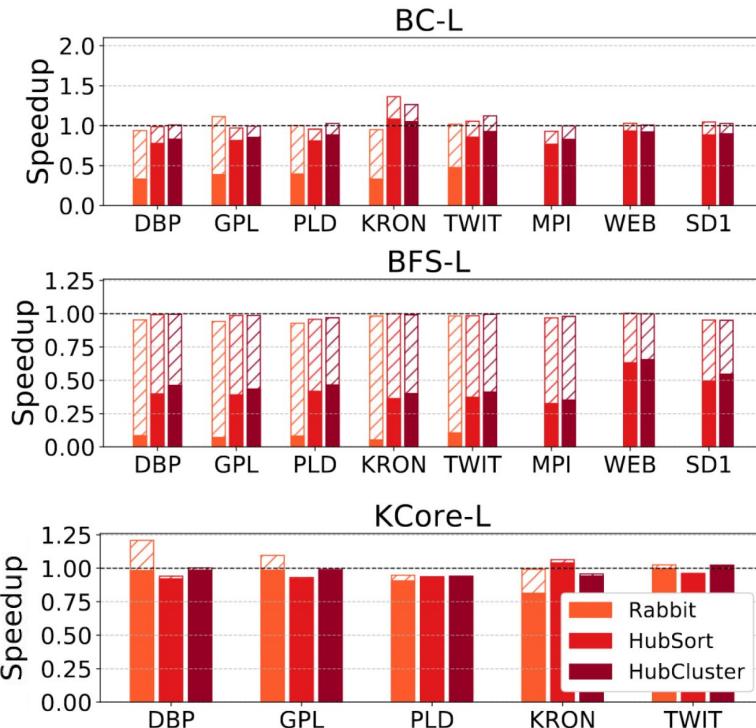
Category III - Applications Processing a Small Fraction of Edges



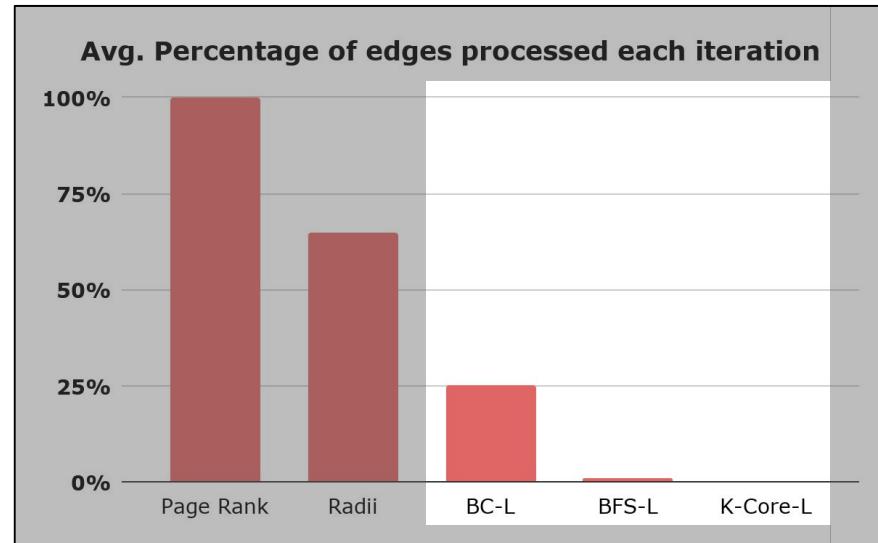
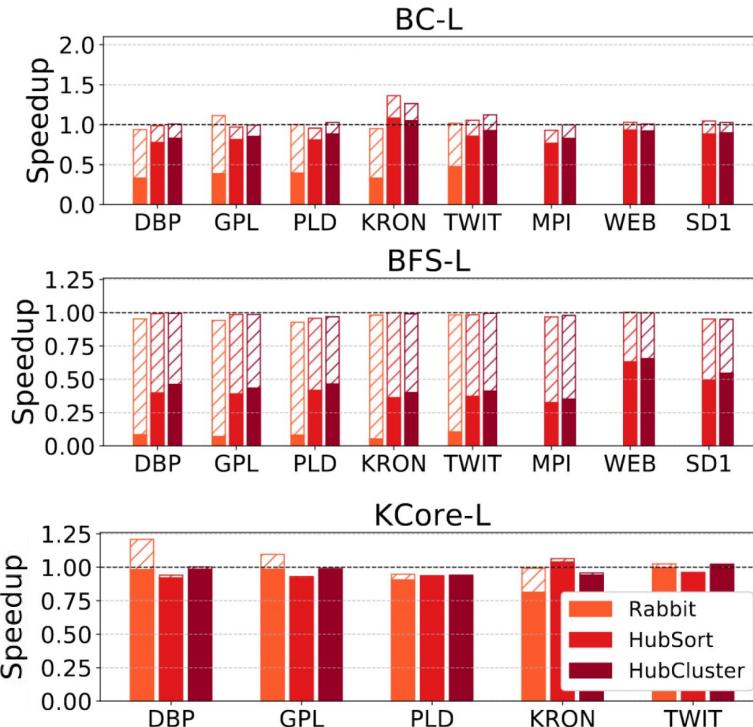
Category III - Applications Processing a Small Fraction of Edges



Category III - Applications Processing a Small Fraction of Edges



Category III - Applications Processing a Small Fraction of Edges



Limited reuse in vtxData accesses
 ↓
 Lower headroom for reordering

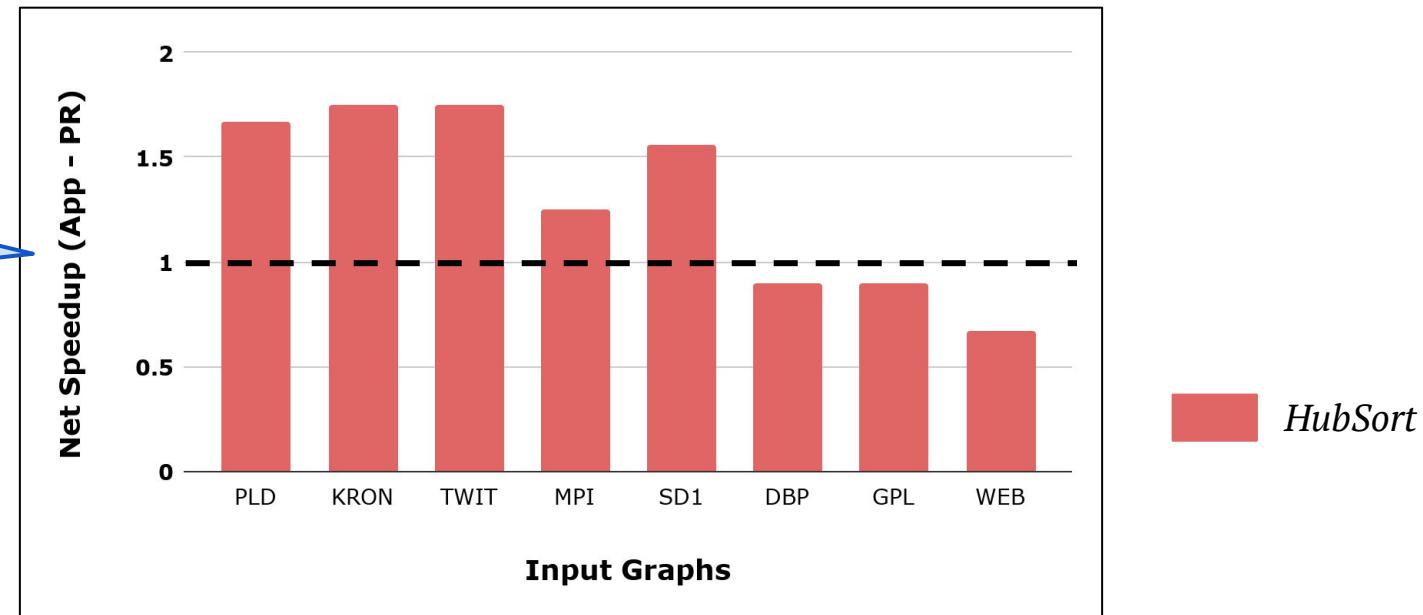
Outline

- ❖ Poor Locality of Graph Processing Applications ✓
- ❖ Improving locality through Graph Reordering ✓
- ❖ Graph Reordering Challenge - *Application and Input-dependent Speedups* ✓
- ❖ **When is Graph Reordering an Optimization?**
 - Characterization Space ✓
 - Which Applications benefit from Reordering? ✓
 - Which Input Graphs benefit from Reordering?
- ❖ Selective Graph Reordering

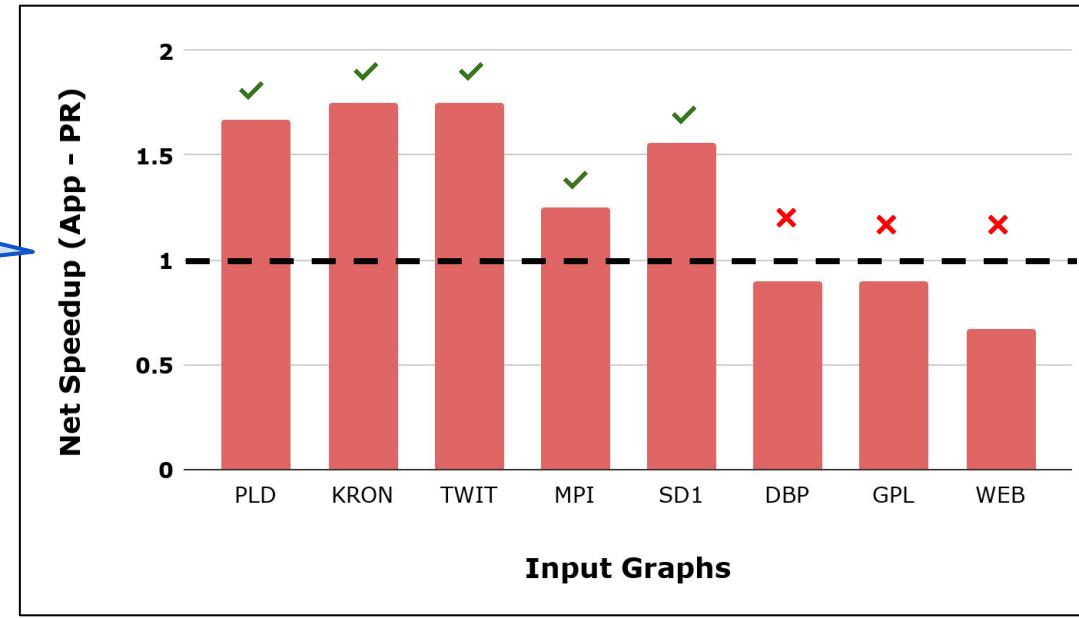
Outline

- ❖ Poor Locality of Graph Processing Applications ✓
- ❖ Improving locality through Graph Reordering ✓
- ❖ Graph Reordering Challenge - *Application and Input-dependent Speedups* ✓
- ❖ **When is Graph Reordering an Optimization?**
 - Characterization Space ✓
 - Which Applications benefit from Reordering? ✓
 - Which Input Graphs benefit from Reordering?
- ❖ Selective Graph Reordering

Speedup From HubSorting Varies Across Inputs



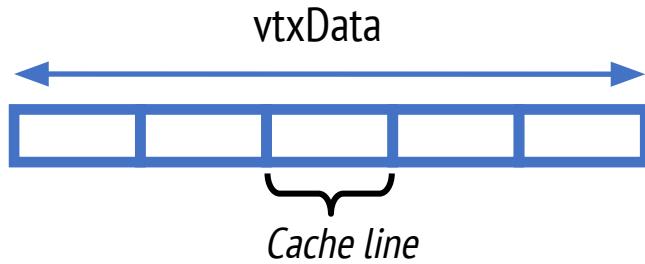
Speedup From HubSorting Varies Across Inputs



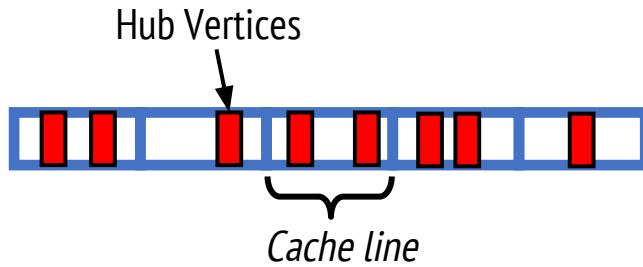
HubSort

Need to *predict speedup* from HubSorting
AND
selectively perform HubSorting

Understanding Performance Improvement From HubSorting

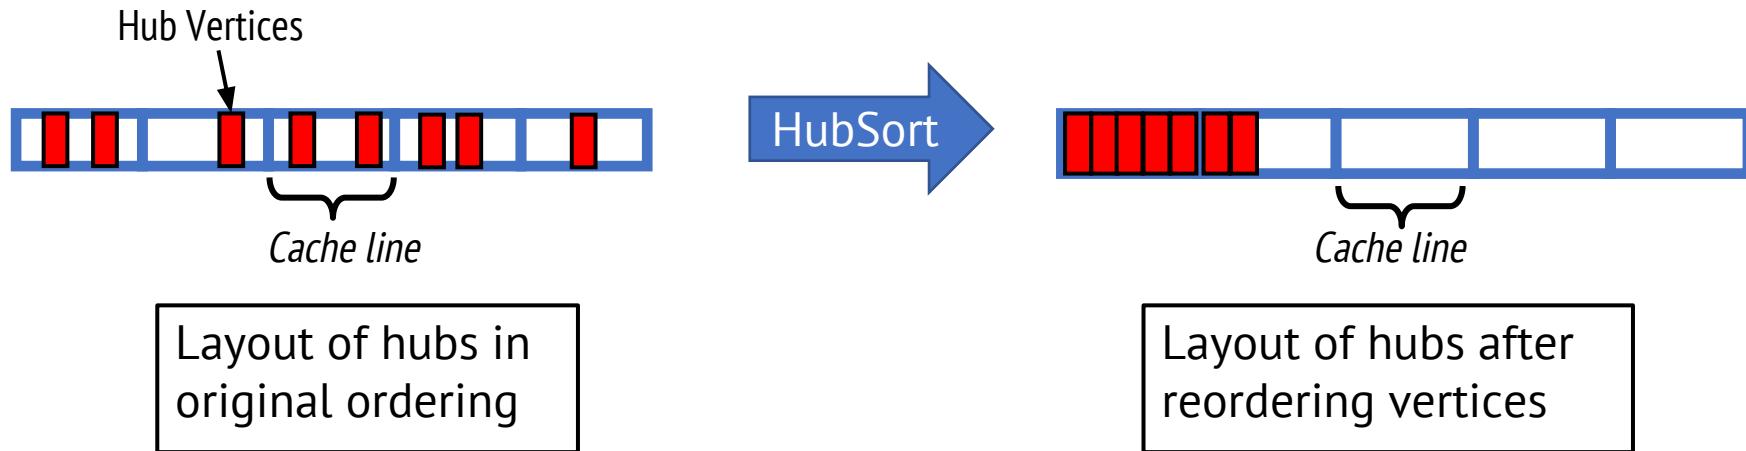


Understanding Performance Improvement From HubSorting

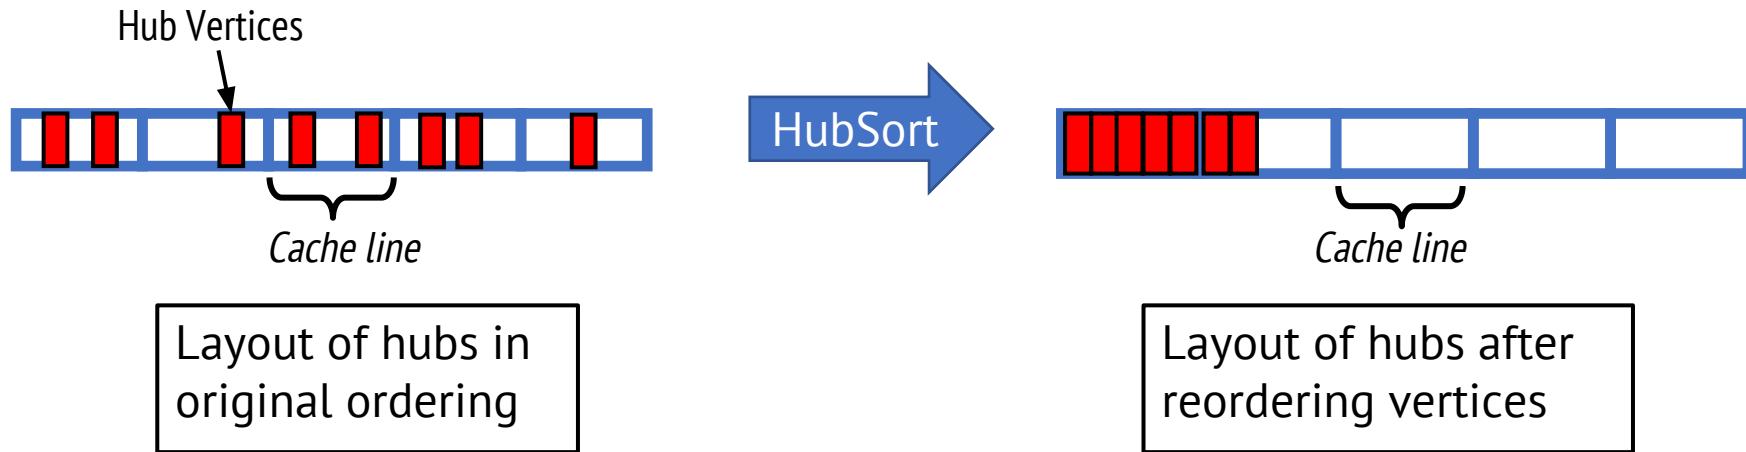


Layout of hubs in
original ordering

Understanding Performance Improvement From HubSorting



Understanding Performance Improvement From HubSorting



HubSorting will be most effective for Graphs with:

- ❖ **Property #1:** Skew in the degree-distribution (*Presence of Hubs*)
- ❖ **Property #2:** Sparsely distributed hub vertices (*Quality of original ordering*)

Packing Factor - A Measure of Hub Density

Packing Factor is a measure of how densely the hubs are packed after HubSorting



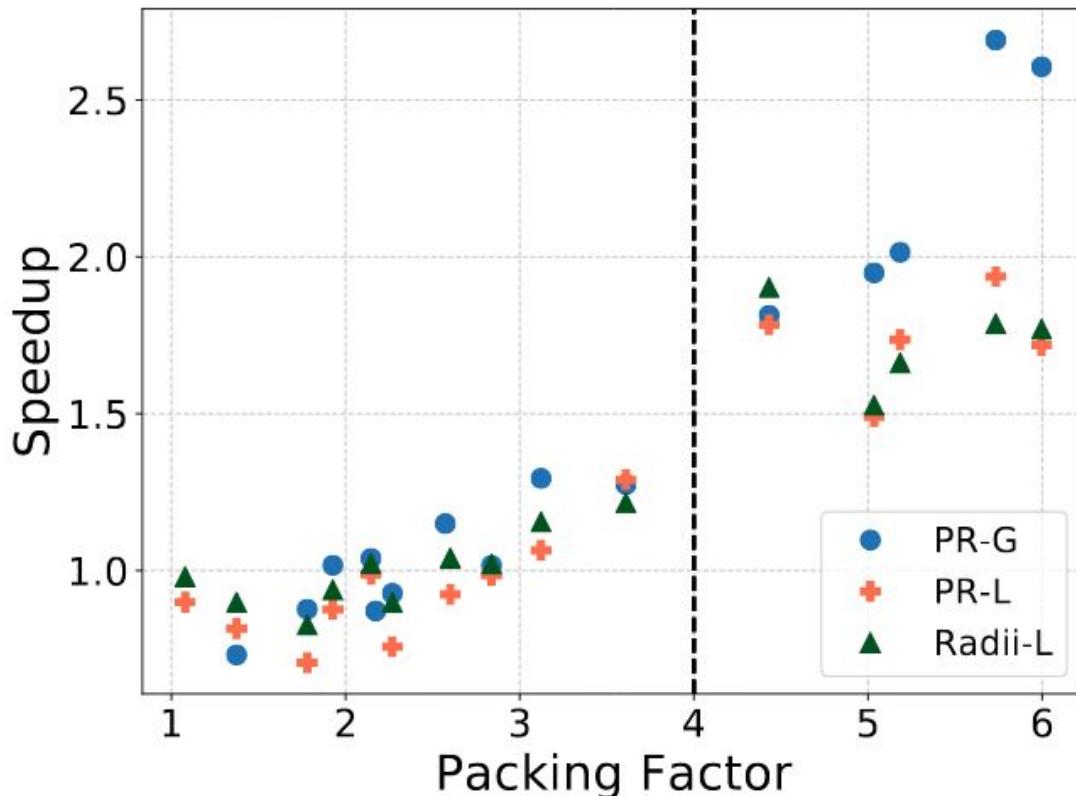
*Layout of the original
ordering of vertices*

*Layout after reordering
vertices by HubSorting*

$$\text{Packing Factor} = \frac{5}{2} = 2.5$$

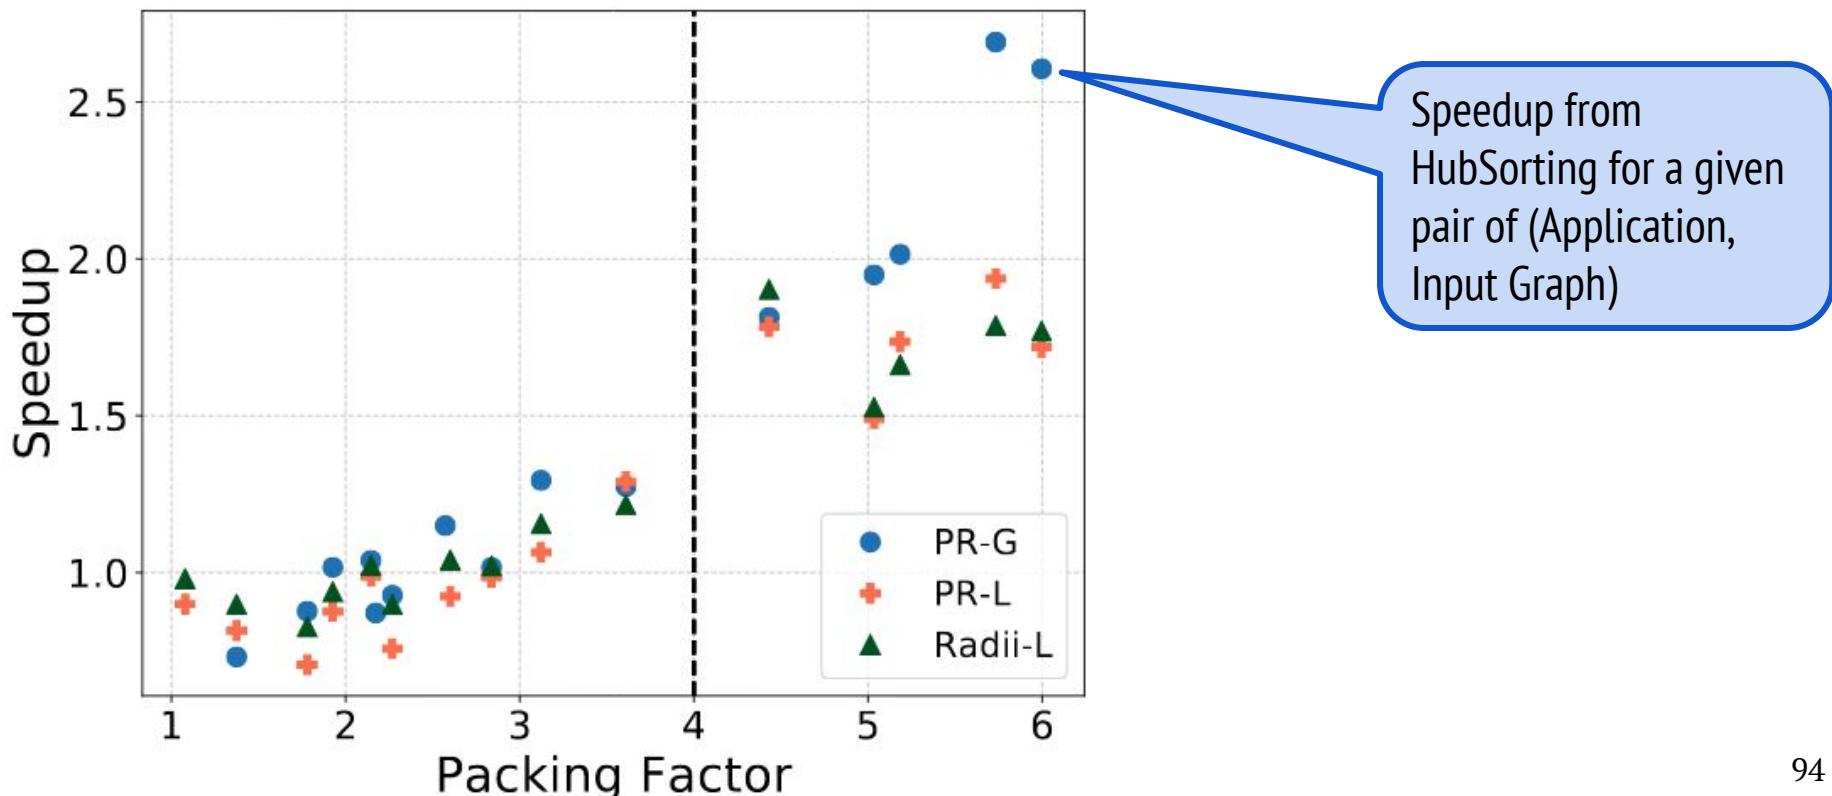
Packing Factor Can Predict Speedup From HubSorting

Pearson Correlation = 0.92



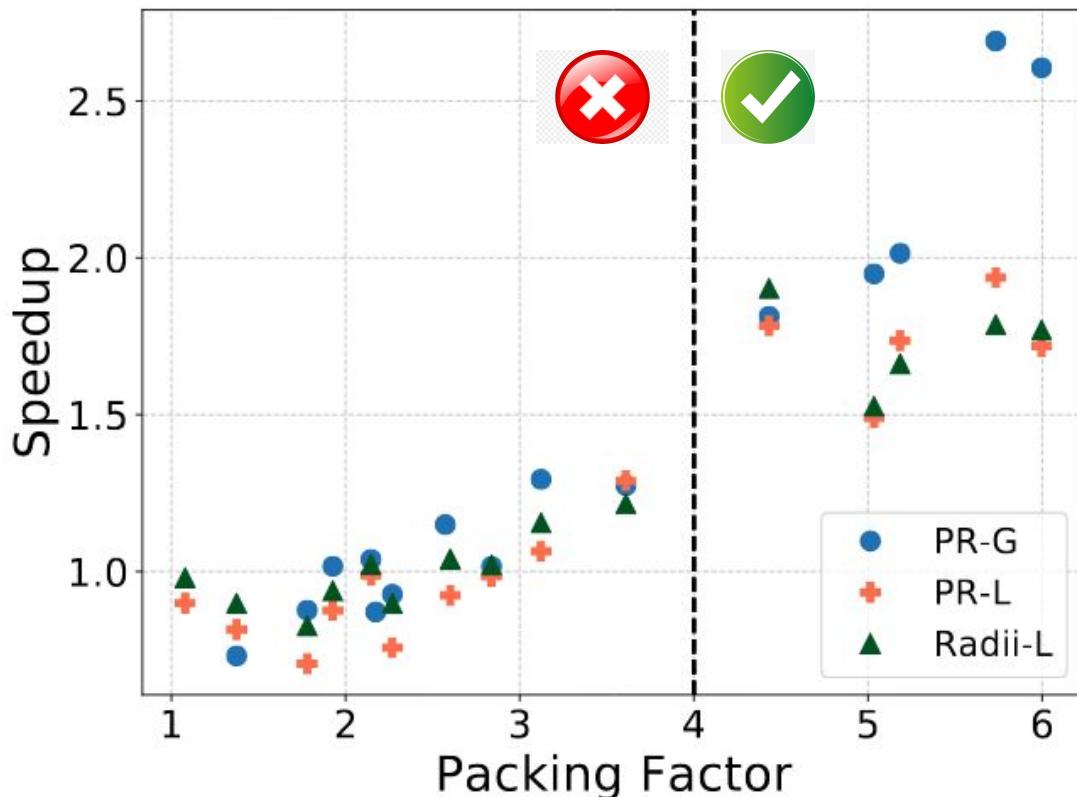
Packing Factor Can Predict Speedup From HubSorting

Pearson Correlation = 0.92



Packing Factor Can Predict Speedup From HubSorting

Pearson Correlation = 0.92



Packing Factor is a good predictor for Speedup

Outline

- ❖ Poor Locality of Graph Processing Applications ✓
- ❖ Improving locality through Graph Reordering ✓
- ❖ Graph Reordering Challenge - *Application and Input-dependent Speedups* ✓
- ❖ When is Graph Reordering an Optimization?
 - Characterization Space ✓
 - Which Applications benefit from Reordering? ✓
 - Which Input Graphs benefit from Reordering? ✓
- ❖ Selective Graph Reordering

Outline

- ❖ Poor Locality of Graph Processing Applications ✓
- ❖ Improving locality through Graph Reordering ✓
- ❖ Graph Reordering Challenge - *Application and Input-dependent Speedups* ✓
- ❖ When is Graph Reordering an Optimization?
 - Characterization Space ✓
 - Which Applications benefit from Reordering? ✓
 - Which Input Graphs benefit from Reordering? ✓
- ❖ **Selective Graph Reordering**

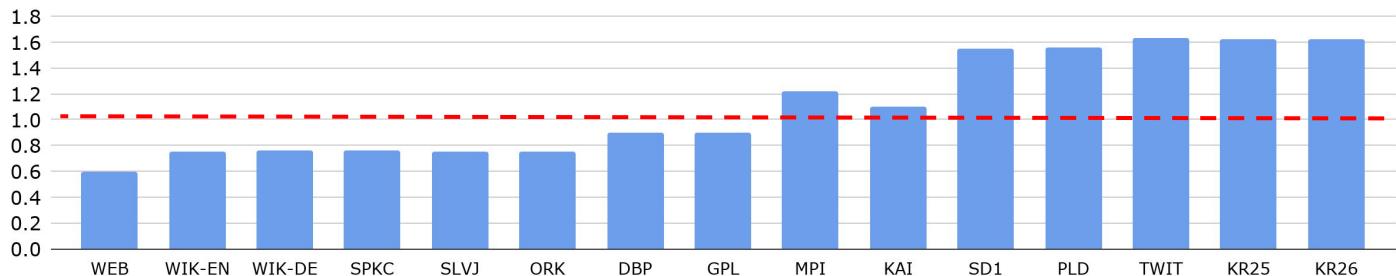
Selective Graph Reordering

```
G` = HubSort(G)
```

```
Process(G`)
```

Selective Graph Reordering

Net Speedup from Unconditionally HubSorting (PR-G)



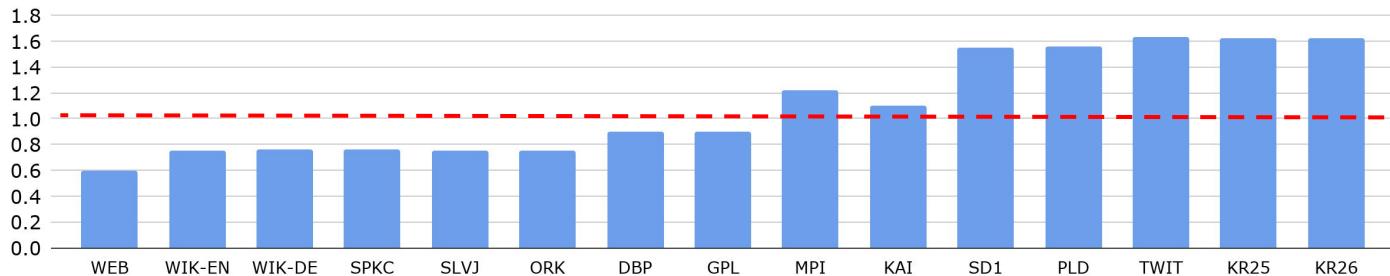
$G' = \text{HubSort}(G)$
Process(G')



Increasing order of Packing Factor

Selective Graph Reordering

Net Speedup from Unconditionally HubSorting (PR-G)

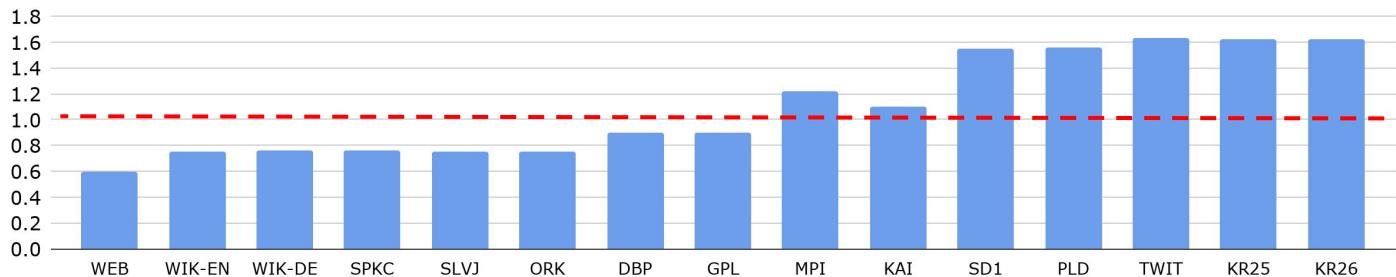


```
G` = HubSort(G)  
Process(G`)
```

```
PF = ComputePF(G)  
if (PF > 4):  
    G` = HubSort(G)  
    Process(G`)  
else:  
    Process(G)
```

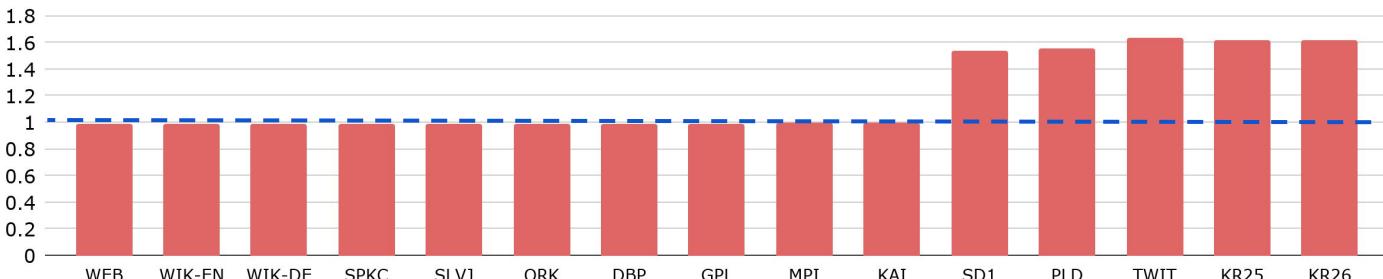
Selective Graph Reordering

Net Speedup from Unconditionally HubSorting (PR-G)



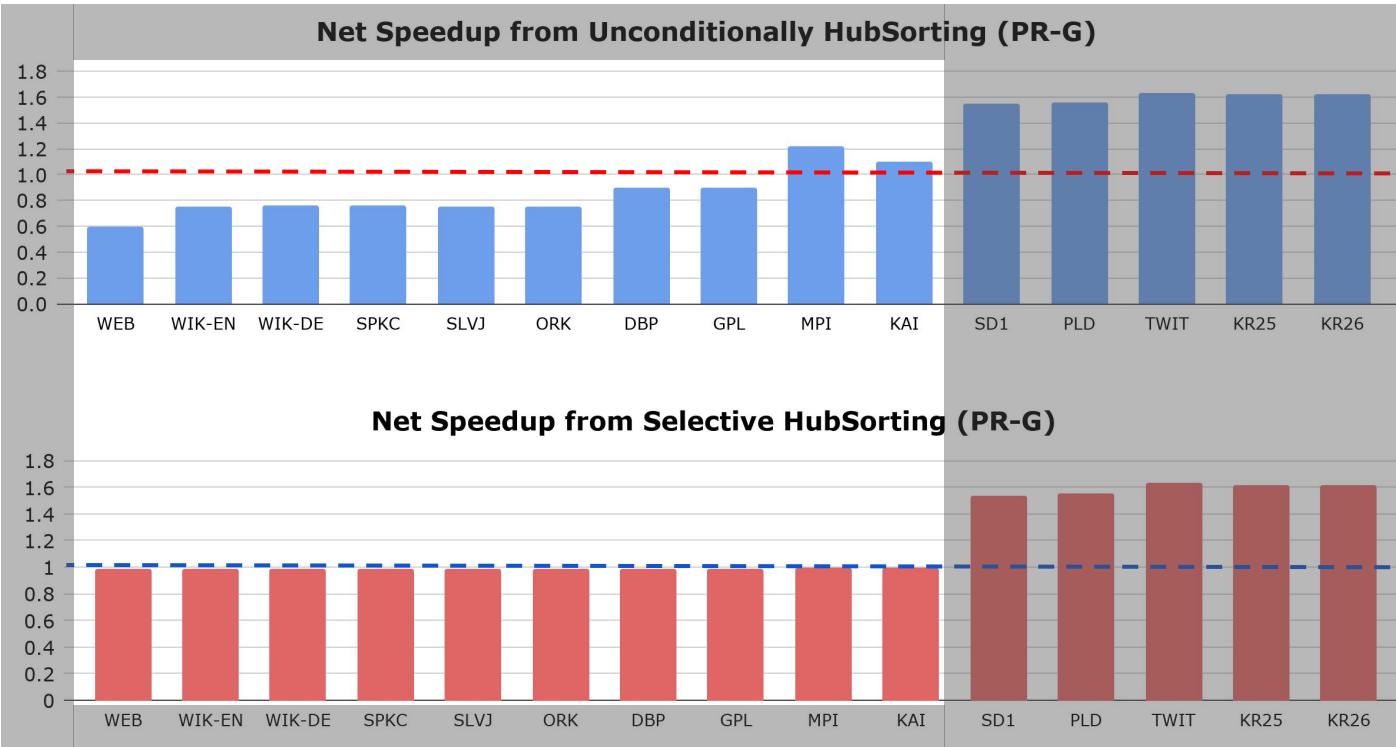
```
G` = HubSort(G)
Process(G`)
```

Net Speedup from Selective HubSorting (PR-G)



```
PF = ComputePF(G)
if (PF > 4):
    G` = HubSort(G)
    Process(G`)
else:
    Process(G)
```

Selective Graph Reordering



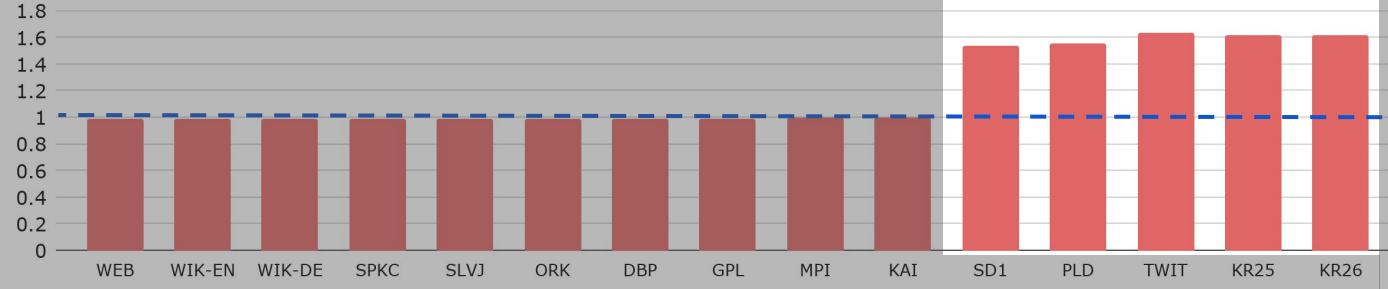
Selective
Reordering
avoids
slowdowns

Selective Graph Reordering

Net Speedup from Unconditionally HubSorting (PR-G)



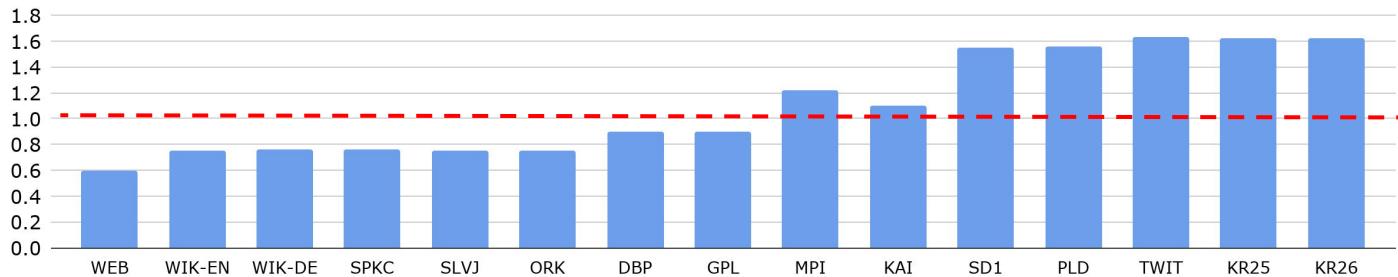
Net Speedup from Selective HubSorting (PR-G)



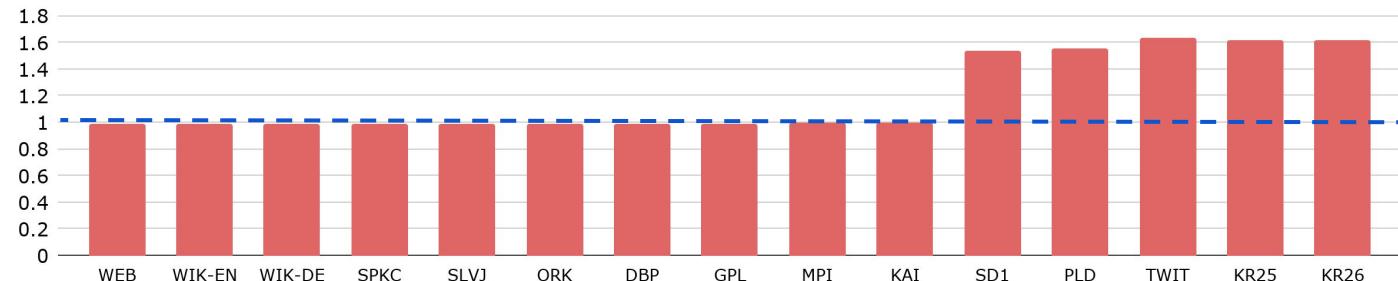
Computing
Packing Factor
does not
degrade
performance

Selective Graph Reordering

Net Speedup from Unconditionally HubSorting (PR-G)



Net Speedup from Selective HubSorting (PR-G)



Selective
Reordering is
a *viable*
Optimization

Conclusions

- ❖ Graph Reordering does not benefit all Application and Input Graphs
- ❖ Opportunity to design new Reordering techniques for specific applications
- ❖ Packing Factor enables Selective Graph Reordering

Source Code Available

- Includes code for:
 - Packing Factor
 - Lightweight Reordering Techniques
 - Selective HubSorting
- Open sourced at -
 - [*https://github.com/CMUAbstract/Graph-Reordering-IISWC18*](https://github.com/CMUAbstract/Graph-Reordering-IISWC18)

Thank You!

When Is Graph Reordering An Optimization?

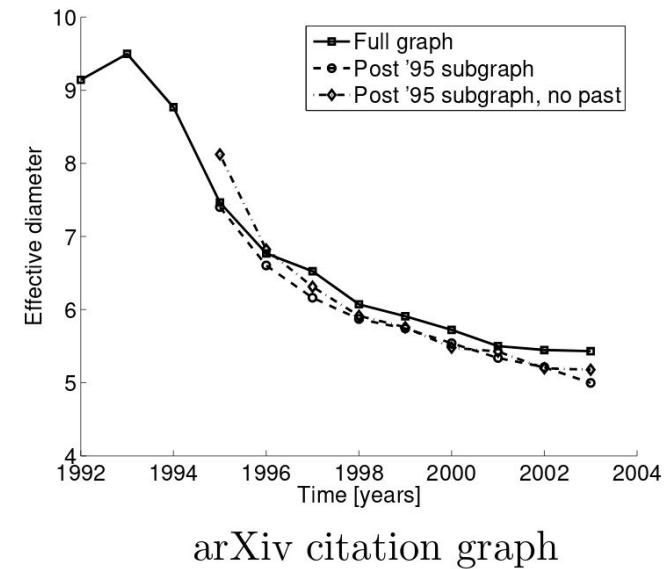
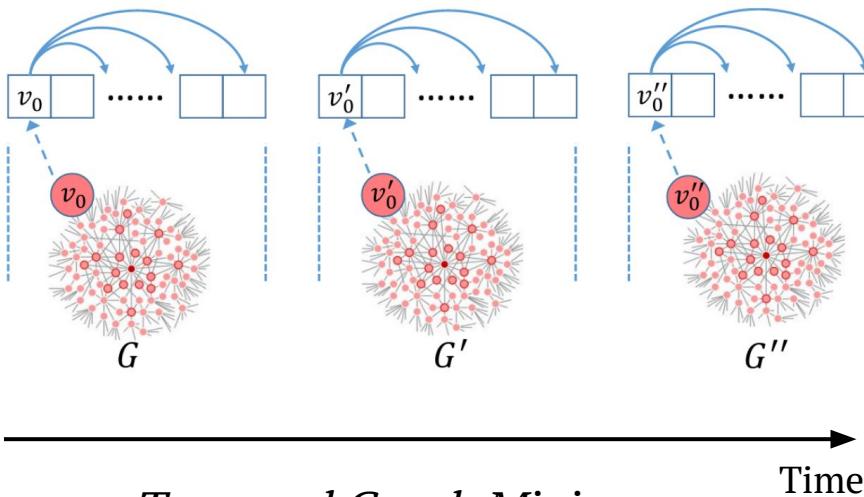
A Cross Application and Input Graph Study on the Effectiveness of Lightweight Graph Reordering

Vignesh Balaji

Brandon Lucia

Backup Slides

Use-cases Where Reordering Overhead Cannot be Amortized



Sophisticated Reordering Techniques are impractical for cases where graph is processed only a few times

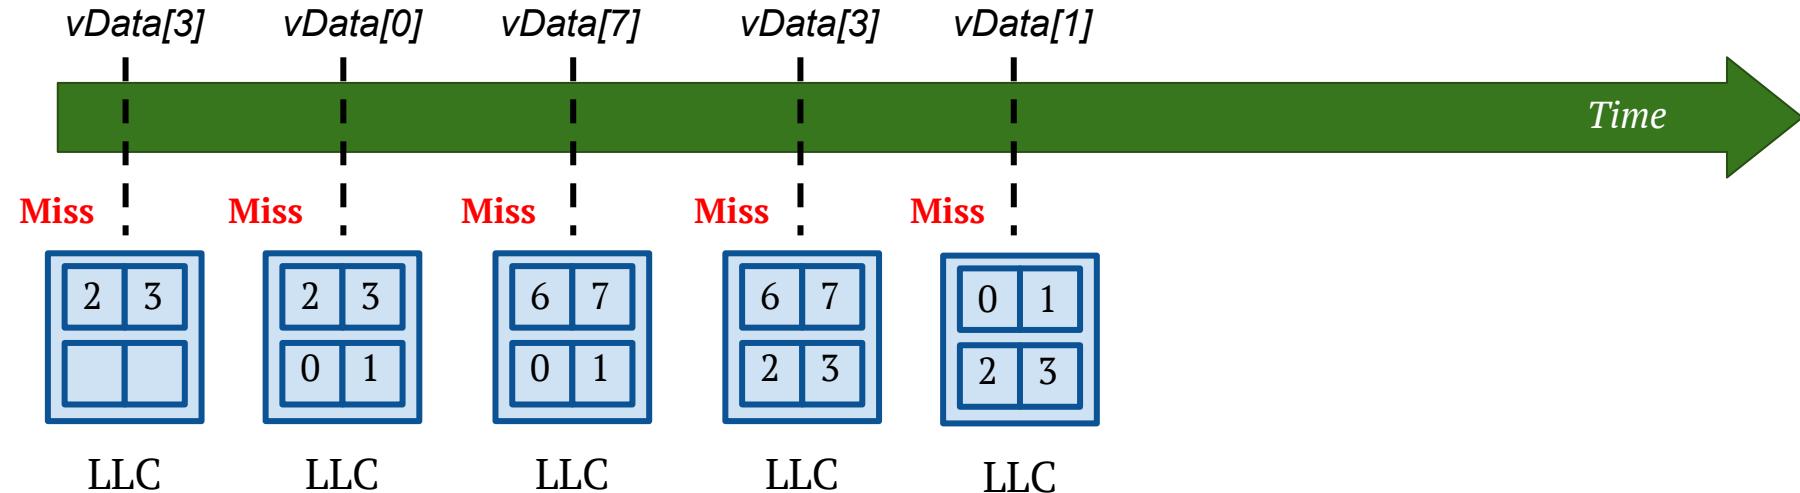
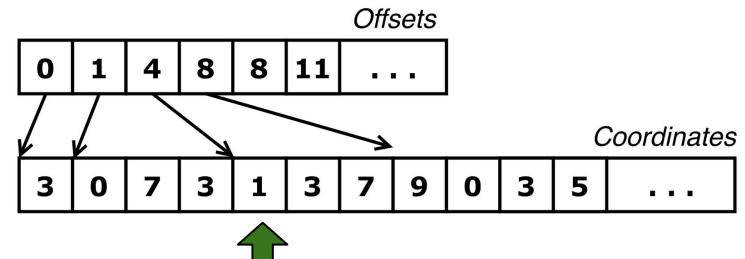
Sophisticated Reordering Techniques Impose *High* Overhead

	gplus	web	pld-arc	twitter	kron26
Run Time (baseline)	6.40s	7.84s	12.40s	21.3s	12.88s
Run Time (Gorder)	4.48s	7.77s	6.54s	13.09s	5.01s
Overhead (Gorder)	1685.9s	459.8s	7255s	25200s	53234s
#Runs to amortize ovhd	873	6477	1237	3072	6771

Assumption: Reordered graph will be processed multiple times

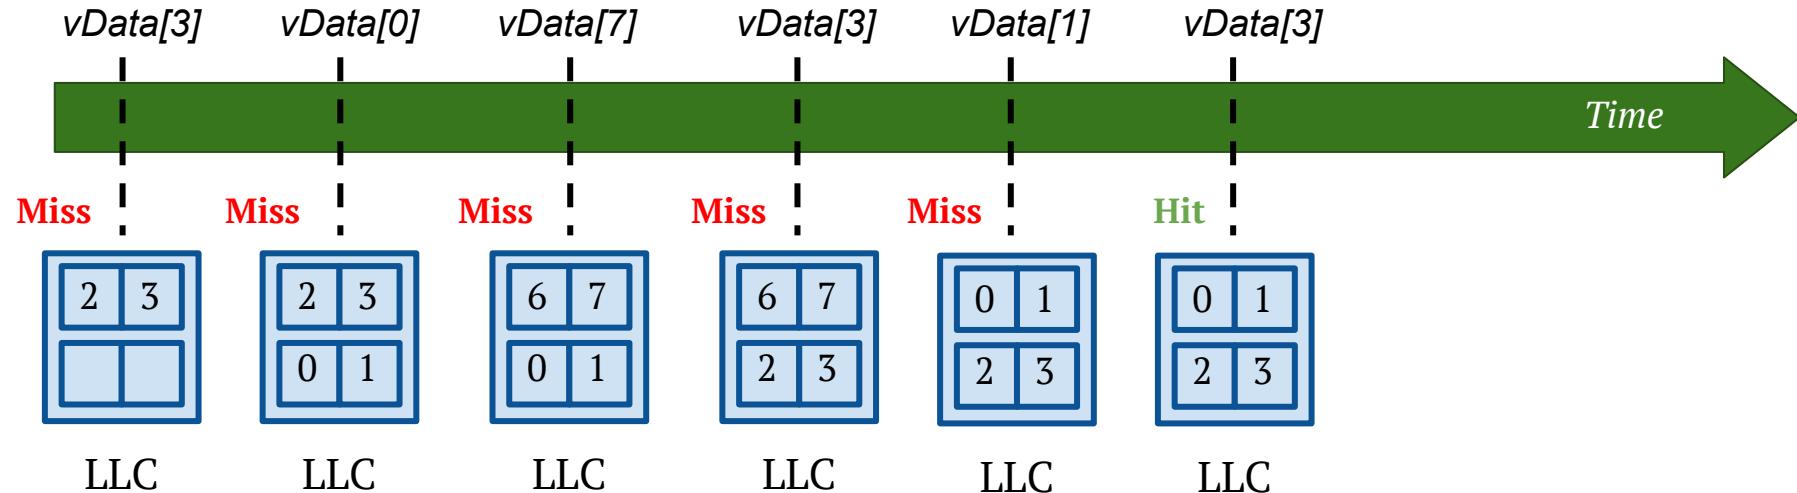
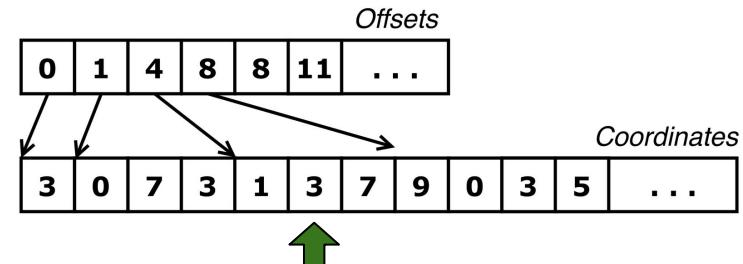
Irregular Accesses Have Poor Temporal And Spatial Locality

```
for v in G:  
    for u in neigh(v):  
        process(..., vData[u], ...)
```



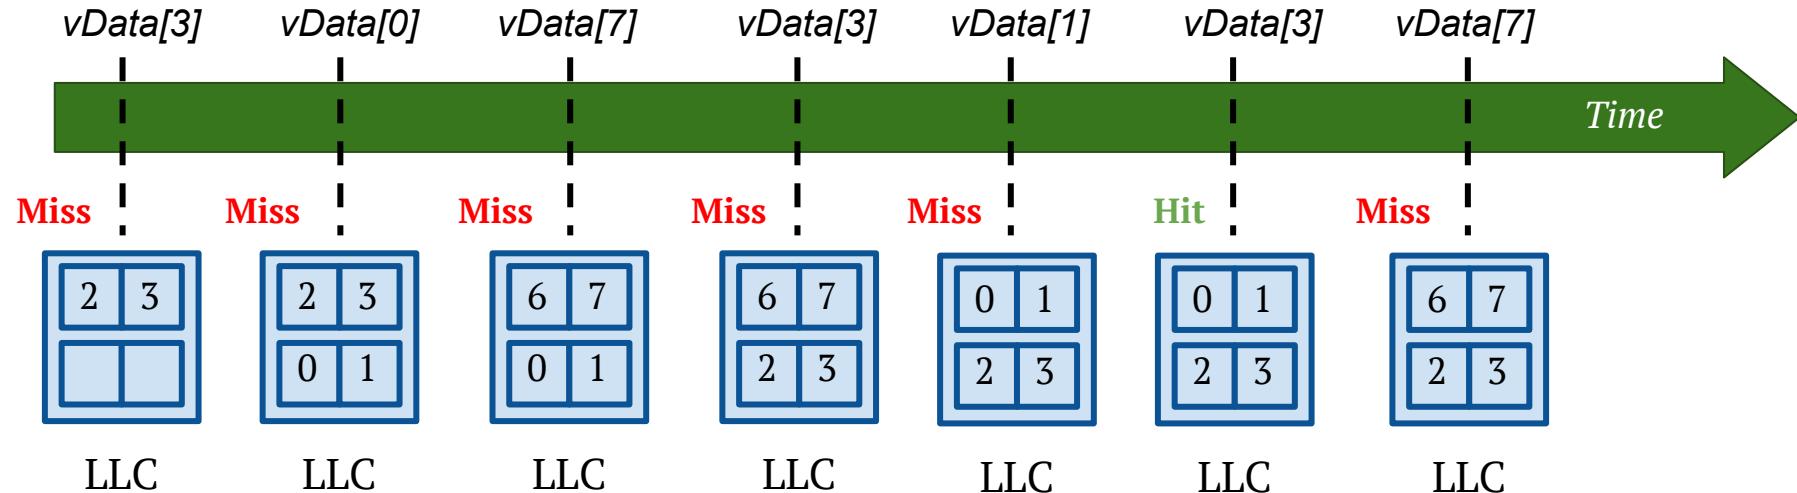
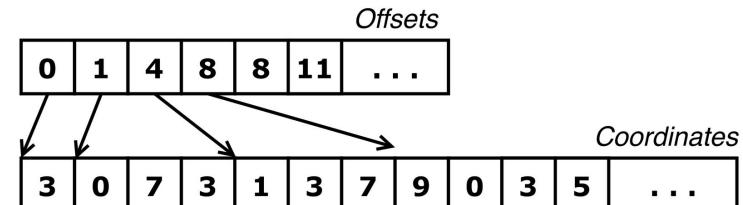
Irregular Accesses Have Poor Temporal And Spatial Locality

```
for v in G:  
    for u in neigh(v):  
        process(..., vData[u], ...)
```



Irregular Accesses Have Poor Temporal And Spatial Locality

```
for v in G:  
    for u in neigh(v):  
        process(..., vData[u], ...)
```



Graph Applications

Ligra

- Page Rank
- Page Rank-Delta
- SSSP – Bellman Ford
- Collaborative Filtering
- Radii
- Betweenness Centrality
- BFS
- Kcore
- Maximal Independent Set
- Connected Components

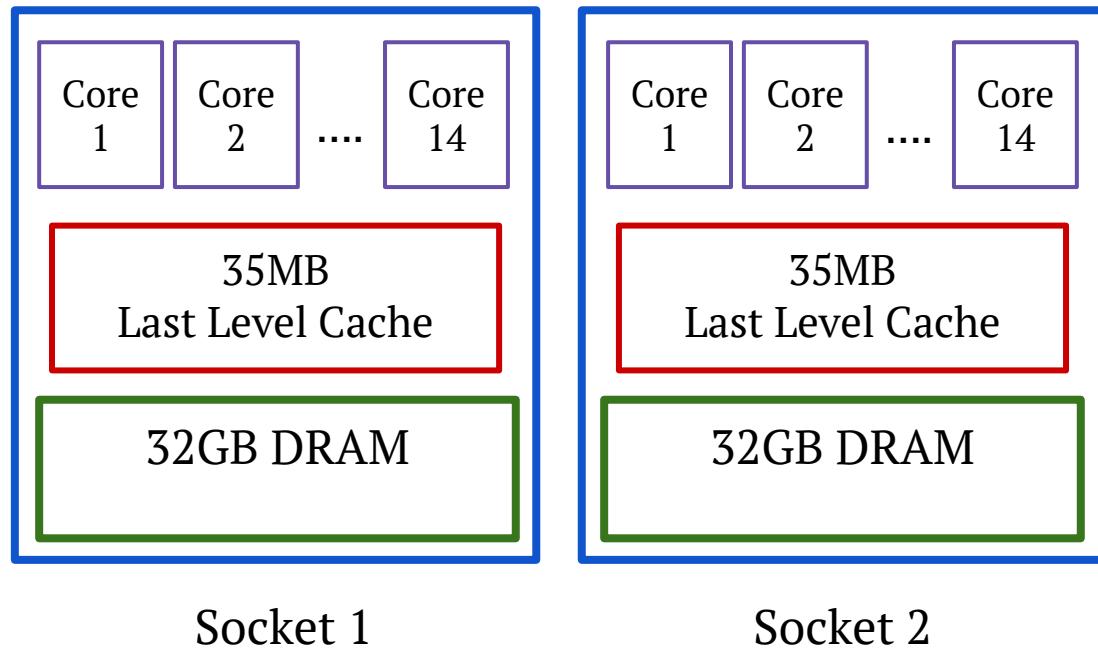
GAP

- Page Rank
- SSSP – Delta Stetting
- Betweenness Centrality
- BFS
- Connected Components

11 Distinct Algorithms

HW Platform

- ❖ Dual-Socket Intel Xeon E5-2660v4 processors
- ❖ 14 cores per Socket (2HT/core)
- ❖ 35 MB Last Level Cache per processor
- ❖ 64 GB of main memory



Input Graphs

	DBP	GPL	PLD	KRON	TWIT	MPI	WEB	SD1
 V (in M)	18.27	28.94	42.89	33.55	61.58	52.58	50.64	94.95
 E (in B)	0.172	0.462	0.623	1.047	1.468	1.963	1.93	1.937
vData Sz (MB)	146.16	231.52	343.12	268.4	498.64	420.64	405.12	759.6
CSR Sz (GB)	1.41	3.66	4.96	8.05	11.34	15.02	14.75	15.13

Input Graphs

	DBP	GPL	PLD	KRON	TWIT	MPI	WEB	SD1
$ V $ (in M)	18.27	28.94	42.89	33.55	61.58	52.58	50.64	94.95
$ E $ (in B)	0.172	0.462	0.623	1.047	1.468	1.963	1.93	1.937
vData Sz (MB)	146.16	231.52	343.12	268.4	498.64	420.64	405.12	759.6
CSR Sz (GB)	1.41	3.66	4.96	8.05	11.34	15.02	14.75	15.13

Irregular working set size >> Aggregate LLC Capacity

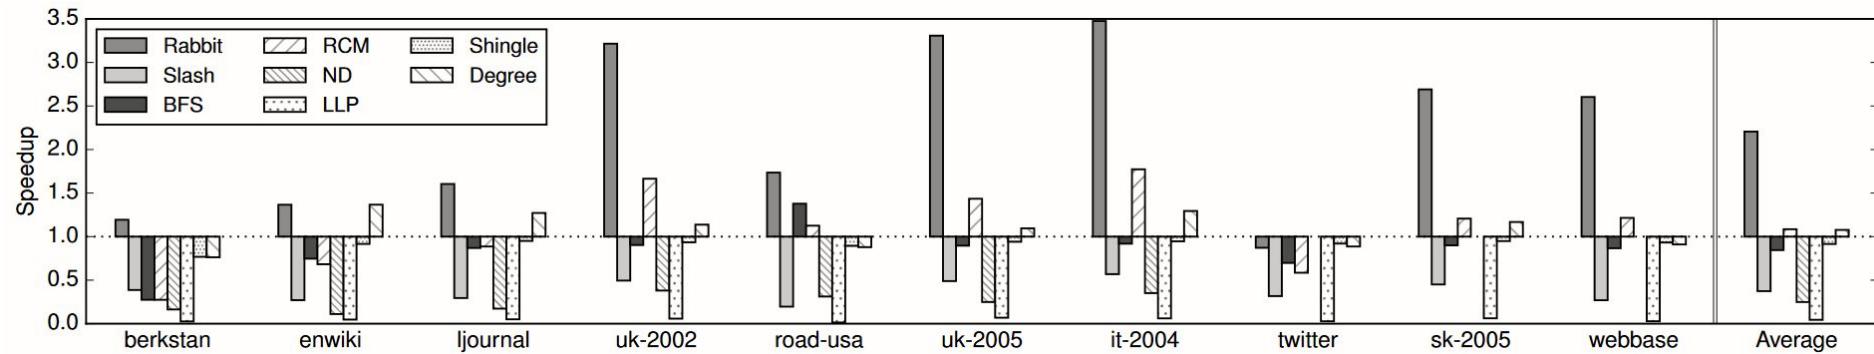
Input Graphs

	DBP	GPL	PLD	KRON	TWIT	MPI	WEB	SD1
$ V $ (in M)	18.27	28.94	42.89	33.55	61.58	52.58	50.64	94.95
$ E $ (in B)	0.172	0.462	0.623	1.047	1.468	1.963	1.93	1.937
vData Sz (MB)	146.16	231.52	343.12	268.4	498.64	420.64	405.12	759.6
CSR Sz (GB)	1.41	3.66	4.96	8.05	11.34	15.02	14.75	15.13

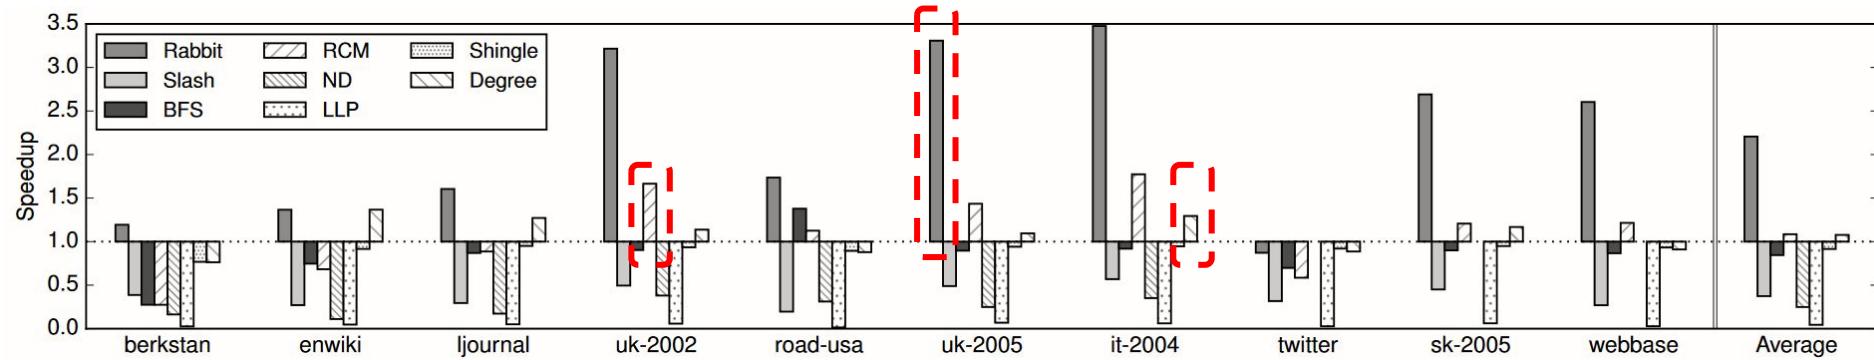
Irregular working set size >> Aggregate LLC Capacity

We use the original ordering of Input Graphs

Lightweight Reordering Can Provide End-to-end Speedups

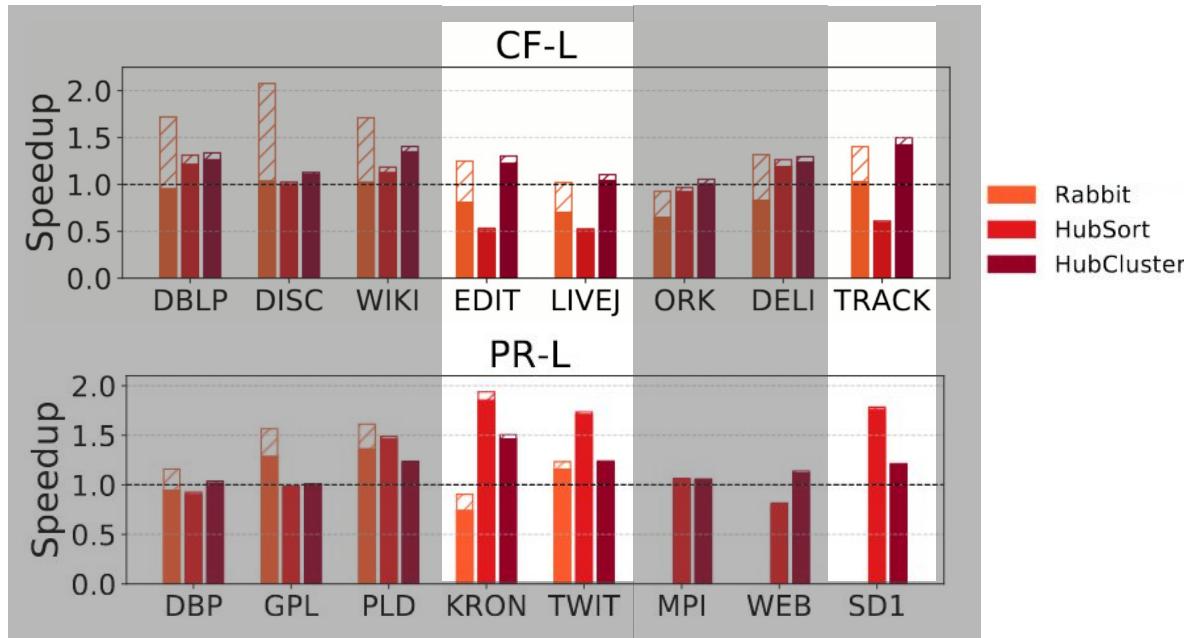


Lightweight Reordering Can Provide End-to-end Speedups



Reordering techniques exploiting
power-law distributions and *community structure*
can have low-overheads

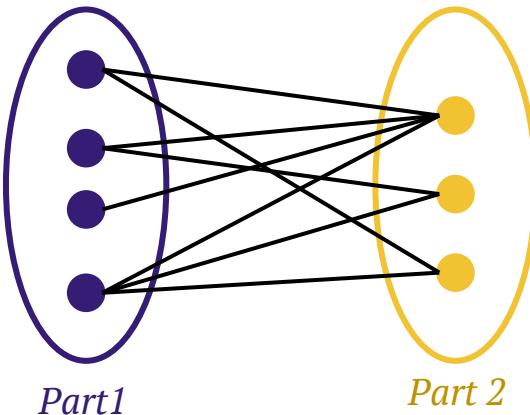
Category II - Executions On Symmetric Bipartite Graphs



Surprising trends:

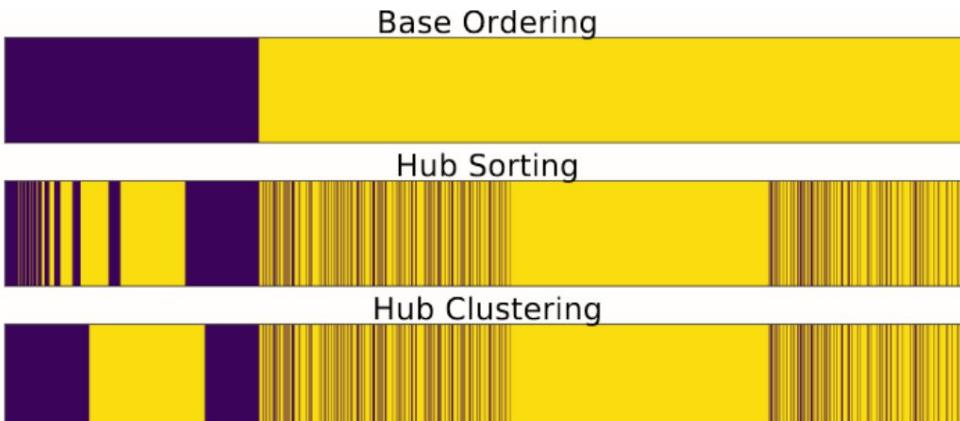
- HubSort offers the least performance benefits
- HubSort causes slowdowns

Category II - Reason For Slowdown With HubSort



```
for v in G:  
    for u in neigh(v):  
        process(..., vData[u], ...)
```

Assigning vertices from each part of the graph a contiguous range is good for temporal locality

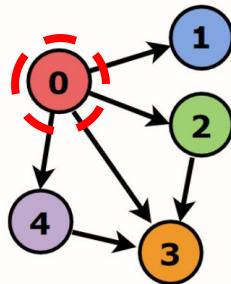


Need a simple mechanism to assign hub vertices from the same part a contiguous range of IDs

Category IV - Push-based Graph Applications

Push-phase

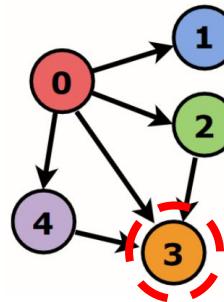
```
parallel for src in Frontier:  
    for dst in outNeigh(v):  
        atomic{parent[dst] = src}
```



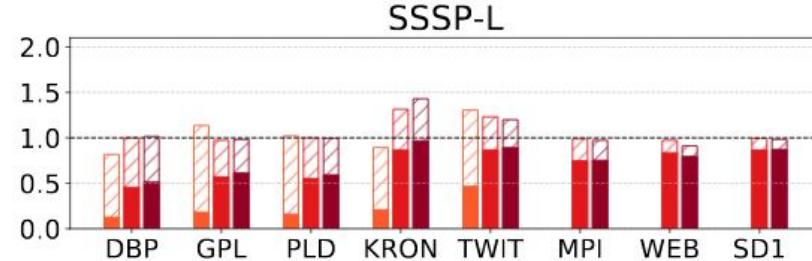
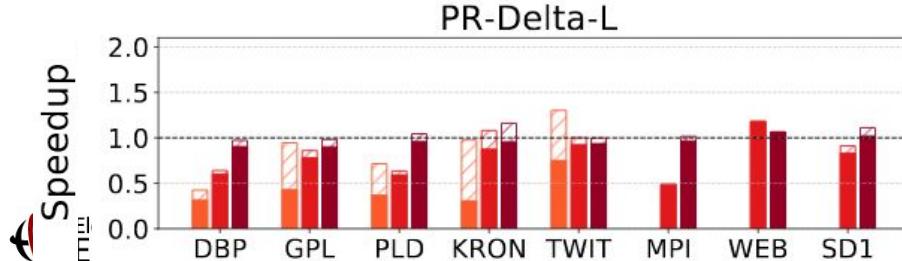
- + Work Efficient execution
- Overhead of synchronization

Pull-phase

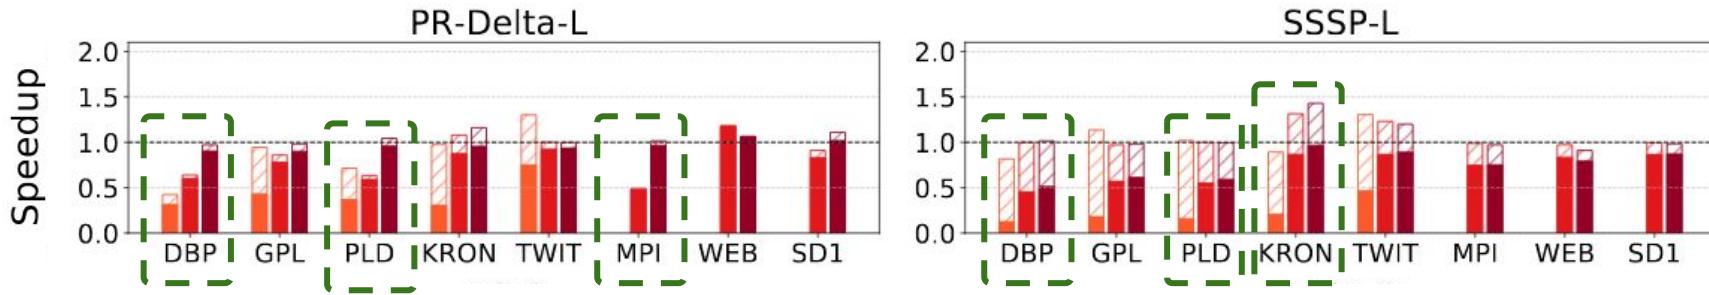
```
parallel for dst in G:  
    for src in inNeigh(v):  
        if src in Frontier:  
            parent[dst] = src
```



- + No synchronization required
- Work-inefficient (iterate over all Vertices)



Category IV - Push-based Graph Applications

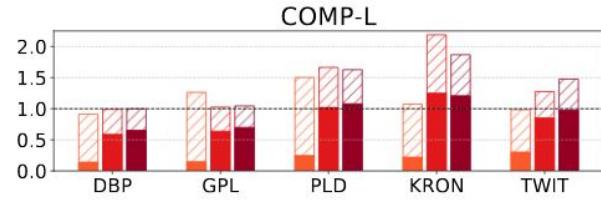
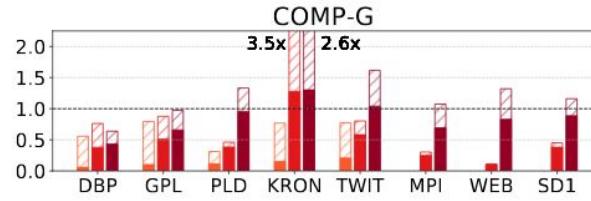
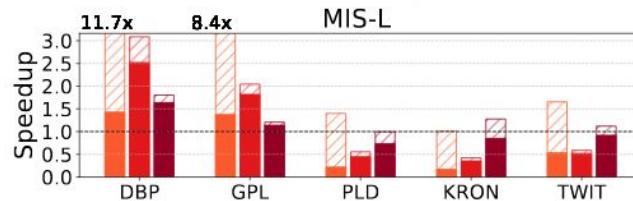


Reordering causes an increase in false-sharing

		DBP	GPL	PLD	KRON	TWIT
PR- δ -L	Rabbit	1.11x	1.53x	1.53x	0.92x	1.26x
	HubSort	0.94x	0.99x	1.43x	1.77x	1.77x
	HubCluster	1.06x	1.01x	1.24x	1.46x	1.27x
SSSP-L	Rabbit	0.87x	1.36x	1.2x	0.95x	0.97x
	HubSort	1.02x	1.14x	1.58x	2.0x	1.4x
	HubCluster	1.14x	1.07x	1.47x	1.58x	1.4x

LWR favors pull-style graph applications

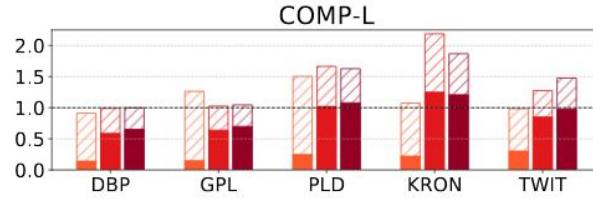
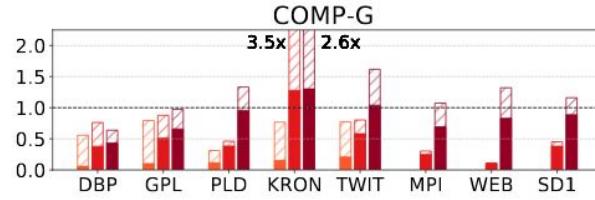
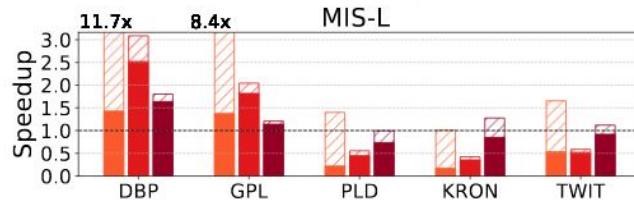
Category V - LWR can affect convergence



```
inline bool update (uintE s, uintE d) {
    //if neighbor is in MIS, then we are out
    if(flags[d] == IN) {if(flags[s] != OUT) flags[s] = OUT;}
    //if neighbor has higher priority (lower ID) and is undecided, then so are we
    else if(d < s)&& flags[s] == CONDITIONALLY_IN && flags[d] < OUT)
        flags[s] = UNDECIDED;
    return 1;
}
```

Vertex IDs influence amount of work done each iteration

Category V - LWR can affect convergence

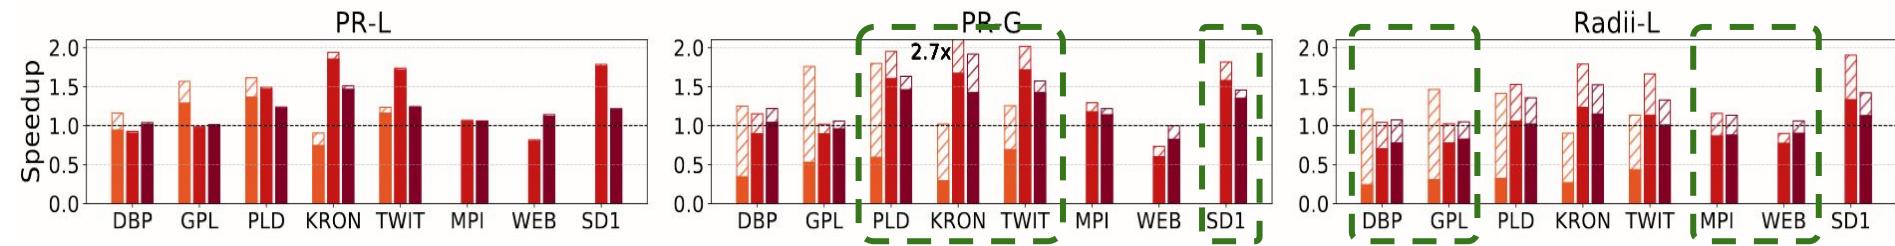


		DBP	GPL	PLD	KRON	TWIT
Comp-G	Rabbit	2.39x	2.0x	5.23x	1.33x	1.47x
	HubSort	1.36x	1.0x	2.09x	0.67x	0.9x
	HubCluster	1.81x	1.0x	1.05x	0.67x	0.88x
Comp-L	Rabbit	1.5x	1.25x	1.27x	1.0x	0.99x
	HubSort	1.25x	1.0x	1.0x	0.67x	0.93x
	HubCluster	1.25x	1.0x	1.0x	0.83x	0.94x
MIS-L	Rabbit	0.3x	0.56x	0.56x	0.96x	0.52x
	HubSort	0.69x	0.56x	0.79x	2.27x	1.01x
	HubCluster	0.85x	0.85x	0.98x	1.19x	1.02x

Opportunity to
accelerate convergence
by reordering vertices

Increase in Iterations until convergence due to LWR

The Need For Selective Lightweight Reordering

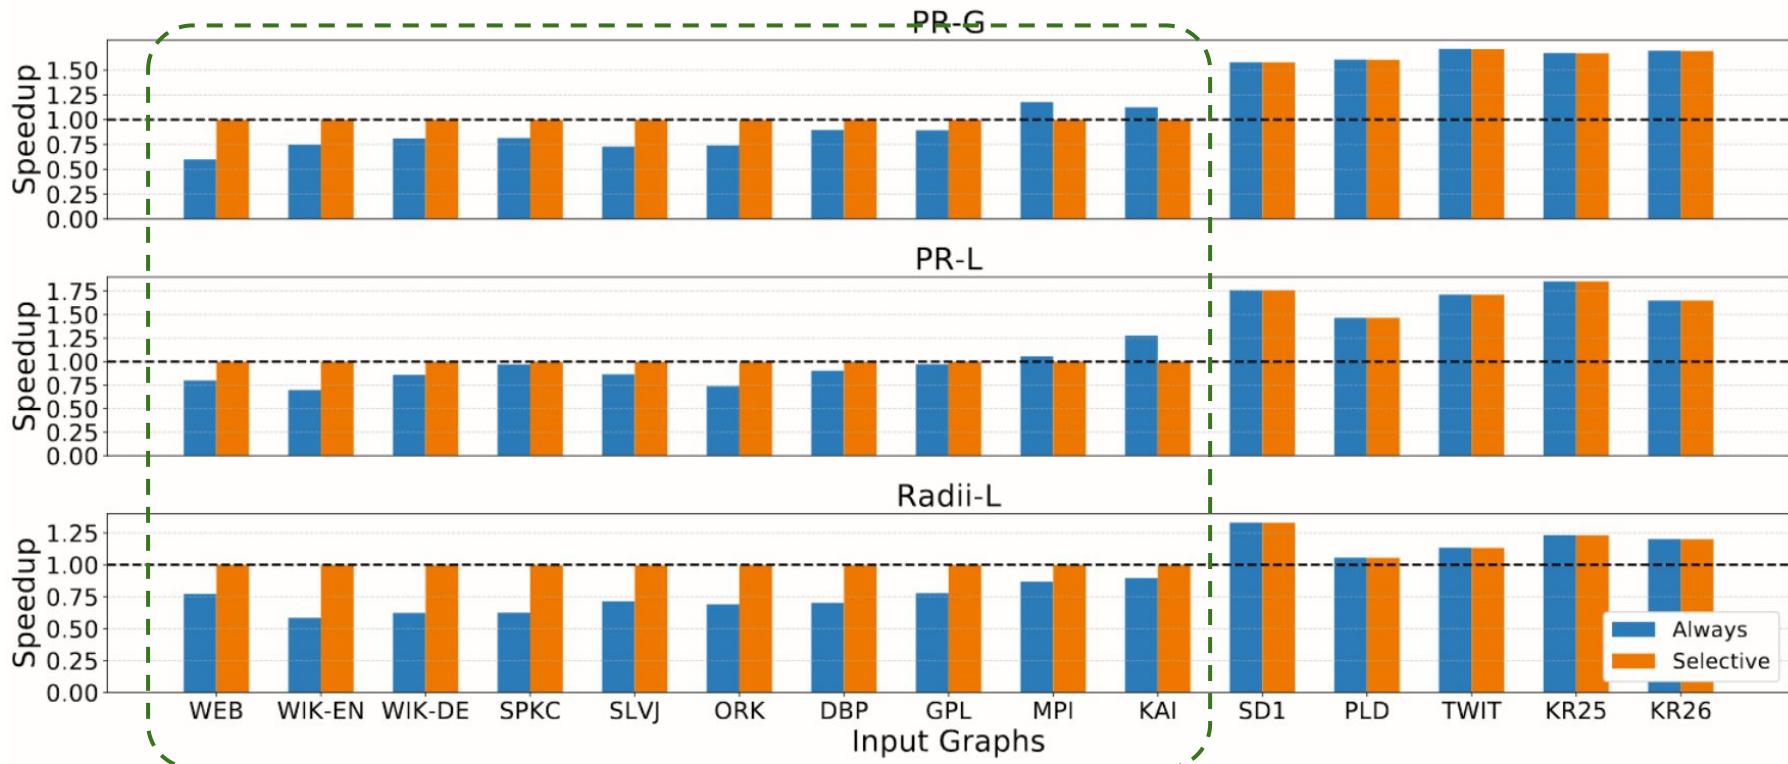


Unconditionally performing LWR causes **net slowdowns** on some input graphs

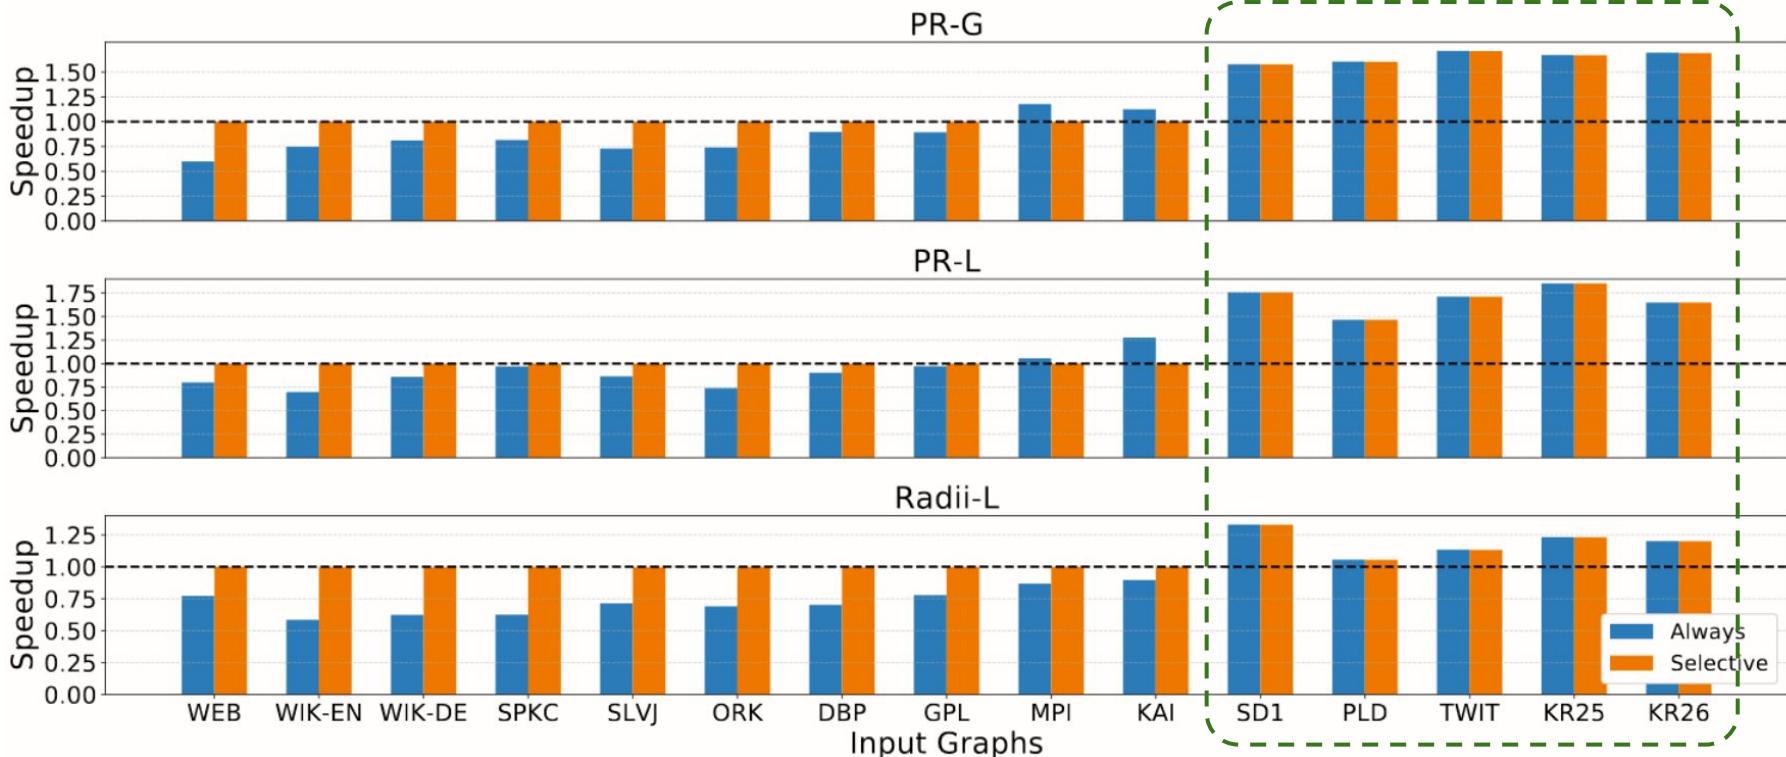
Completely avoid LWR **misses speedups** up to 1.8x

Need to predict speedup from LWR for an input graph and only selectively perform LWR

Using Packing Factor for Selective Reordering

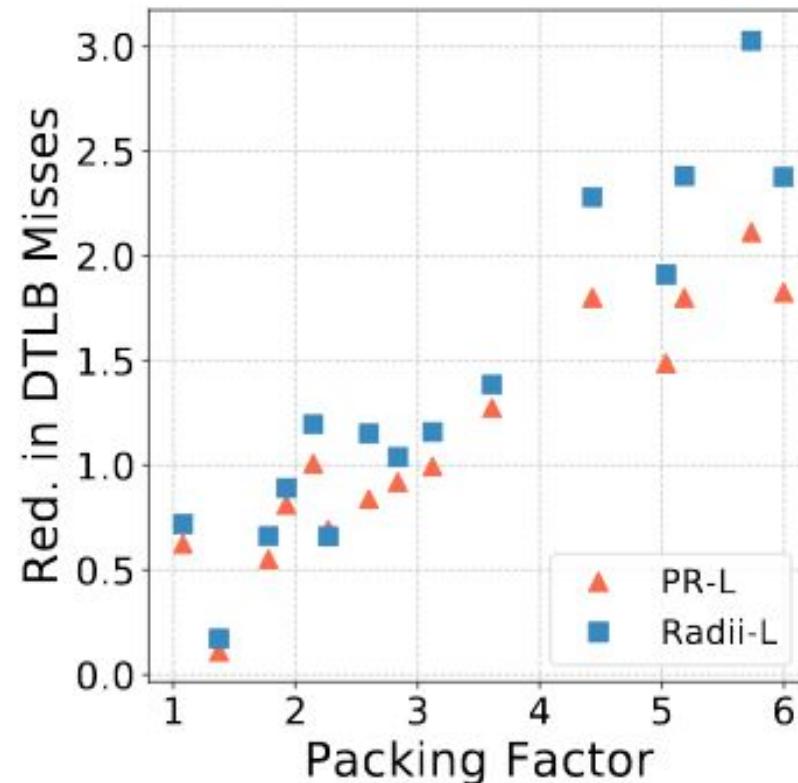
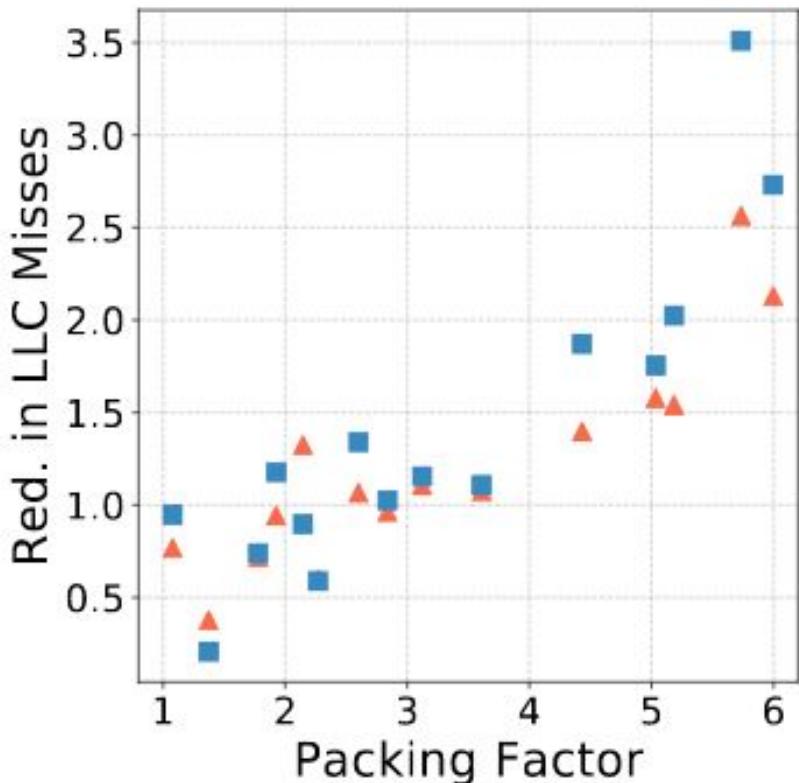


Using Packing Factor for Selective Reordering



The low overhead of Packing Factor computation **does not sacrifice speedup** for high Packing Factor graphs

Speedups From HubSorting Are Due To Locality Improvements



Computing Packing Factor

Algorithm 2 Computing the Packing Factor of a graph

```
1: procedure COMPUTEPACKINGFACTOR( $G$ )
2:    $numHubs \leftarrow 0$ 
3:    $hubWSet\_Original \leftarrow 0$ 
4:   for  $CacheLine$  in  $vDataLines$  do
5:      $containsHub \leftarrow False$ 
6:     for  $vtx$  in  $CacheLine$  do
7:       if ISHUB( $vtx$ ) then
8:          $numHubs += 1$ 
9:          $containsHub \leftarrow True$ 
10:        if  $containsHub = True$  then
11:           $hubWSet\_Original += 1$ 
12:         $hubWSet\_Sorted \leftarrow CEIL(numHubs/VtxPerLine)$ 
13:         $PackingFactor \leftarrow hubWSet\_Original/hubWSet\_Sorted$ 
return  $PackingFactor$ 
```
